

**NUS-ISS**

*Problem Solving Using  
Pattern Recognition*



# Convolutional neural network

by Nicholas Ho

**Recap:**  
**Convolution = filtering**

# 2D convolution

The original

|   | A | B            | C | D | E | F | G | H | I             | J | K | L | M | N | O             | P  | Q |
|---|---|--------------|---|---|---|---|---|---|---------------|---|---|---|---|---|---------------|----|---|
| 1 |   |              |   |   |   |   |   |   |               |   |   |   |   |   |               |    |   |
| 2 |   | <u>Input</u> |   |   |   |   |   |   | <u>Kernel</u> |   |   |   |   |   | <u>Output</u> |    |   |
| 3 |   |              |   |   |   |   |   |   |               |   |   |   |   |   |               |    |   |
| 4 |   | 1            | 3 | 2 | 1 |   |   |   | 1             | 2 | 3 |   |   |   |               |    |   |
| 5 |   | 1            | 3 | 3 | 1 |   |   |   | 0             | 1 | 0 |   |   |   | 23            | 22 |   |
| 6 |   | 2            | 1 | 1 | 3 |   |   |   | 2             | 1 | 2 |   |   |   | 31            | 26 |   |
| 7 |   | 3            | 2 | 3 | 3 |   |   |   |               |   |   |   |   |   |               |    |   |
| 8 |   |              |   |   |   |   |   |   |               |   |   |   |   |   |               |    |   |

## Note:

- Conv1D is used for input signals which are similar to the voice
- Conv2D is used for images
- Conv3D is usually used for videos where you have a frame for each time span

Source: <https://medium.com/apache-mxnet/convolutions-explained-with-ms-excel-465d6649831c>

# 2D convolution

The padded

**Note that the Kernel's movement is determined by the stride value, which can be adjusted (in the examples stride = 1 or [1, 1])**

|    | A | B | C | D            | E | F | G | H | I | J | K             | L | M | N | O  | P  | Q             | R  | S |
|----|---|---|---|--------------|---|---|---|---|---|---|---------------|---|---|---|----|----|---------------|----|---|
| 1  |   |   |   |              |   |   |   |   |   |   |               |   |   |   |    |    |               |    |   |
| 2  |   |   |   | <u>Input</u> |   |   |   |   |   |   | <u>Kernel</u> |   |   |   |    |    | <u>Output</u> |    |   |
| 3  |   |   |   |              |   |   |   |   |   |   |               |   |   |   |    |    |               |    |   |
| 4  |   | 0 | 0 | 0            | 0 | 0 | 0 |   |   |   |               |   |   |   |    |    |               |    |   |
| 5  |   | 0 | 1 | 3            | 2 | 1 | 0 |   |   | 1 | 2             | 3 |   |   | 8  | 14 | 13            | 8  |   |
| 6  |   | 0 | 1 | 3            | 3 | 1 | 0 |   |   | 0 | 1             | 0 |   |   | 16 | 23 | 22            | 10 |   |
| 7  |   | 0 | 2 | 1            | 1 | 3 | 0 |   |   | 2 | 1             | 2 |   |   | 20 | 31 | 26            | 17 |   |
| 8  |   | 0 | 3 | 2            | 3 | 3 | 0 |   |   |   |               |   |   |   | 10 | 9  | 15            | 10 |   |
| 9  |   | 0 | 0 | 0            | 0 | 0 | 0 |   |   |   |               |   |   |   |    |    |               |    |   |
| 10 |   |   |   |              |   |   |   |   |   |   |               |   |   |   |    |    |               |    |   |

**Padding retains the size of the output; input size = output size**

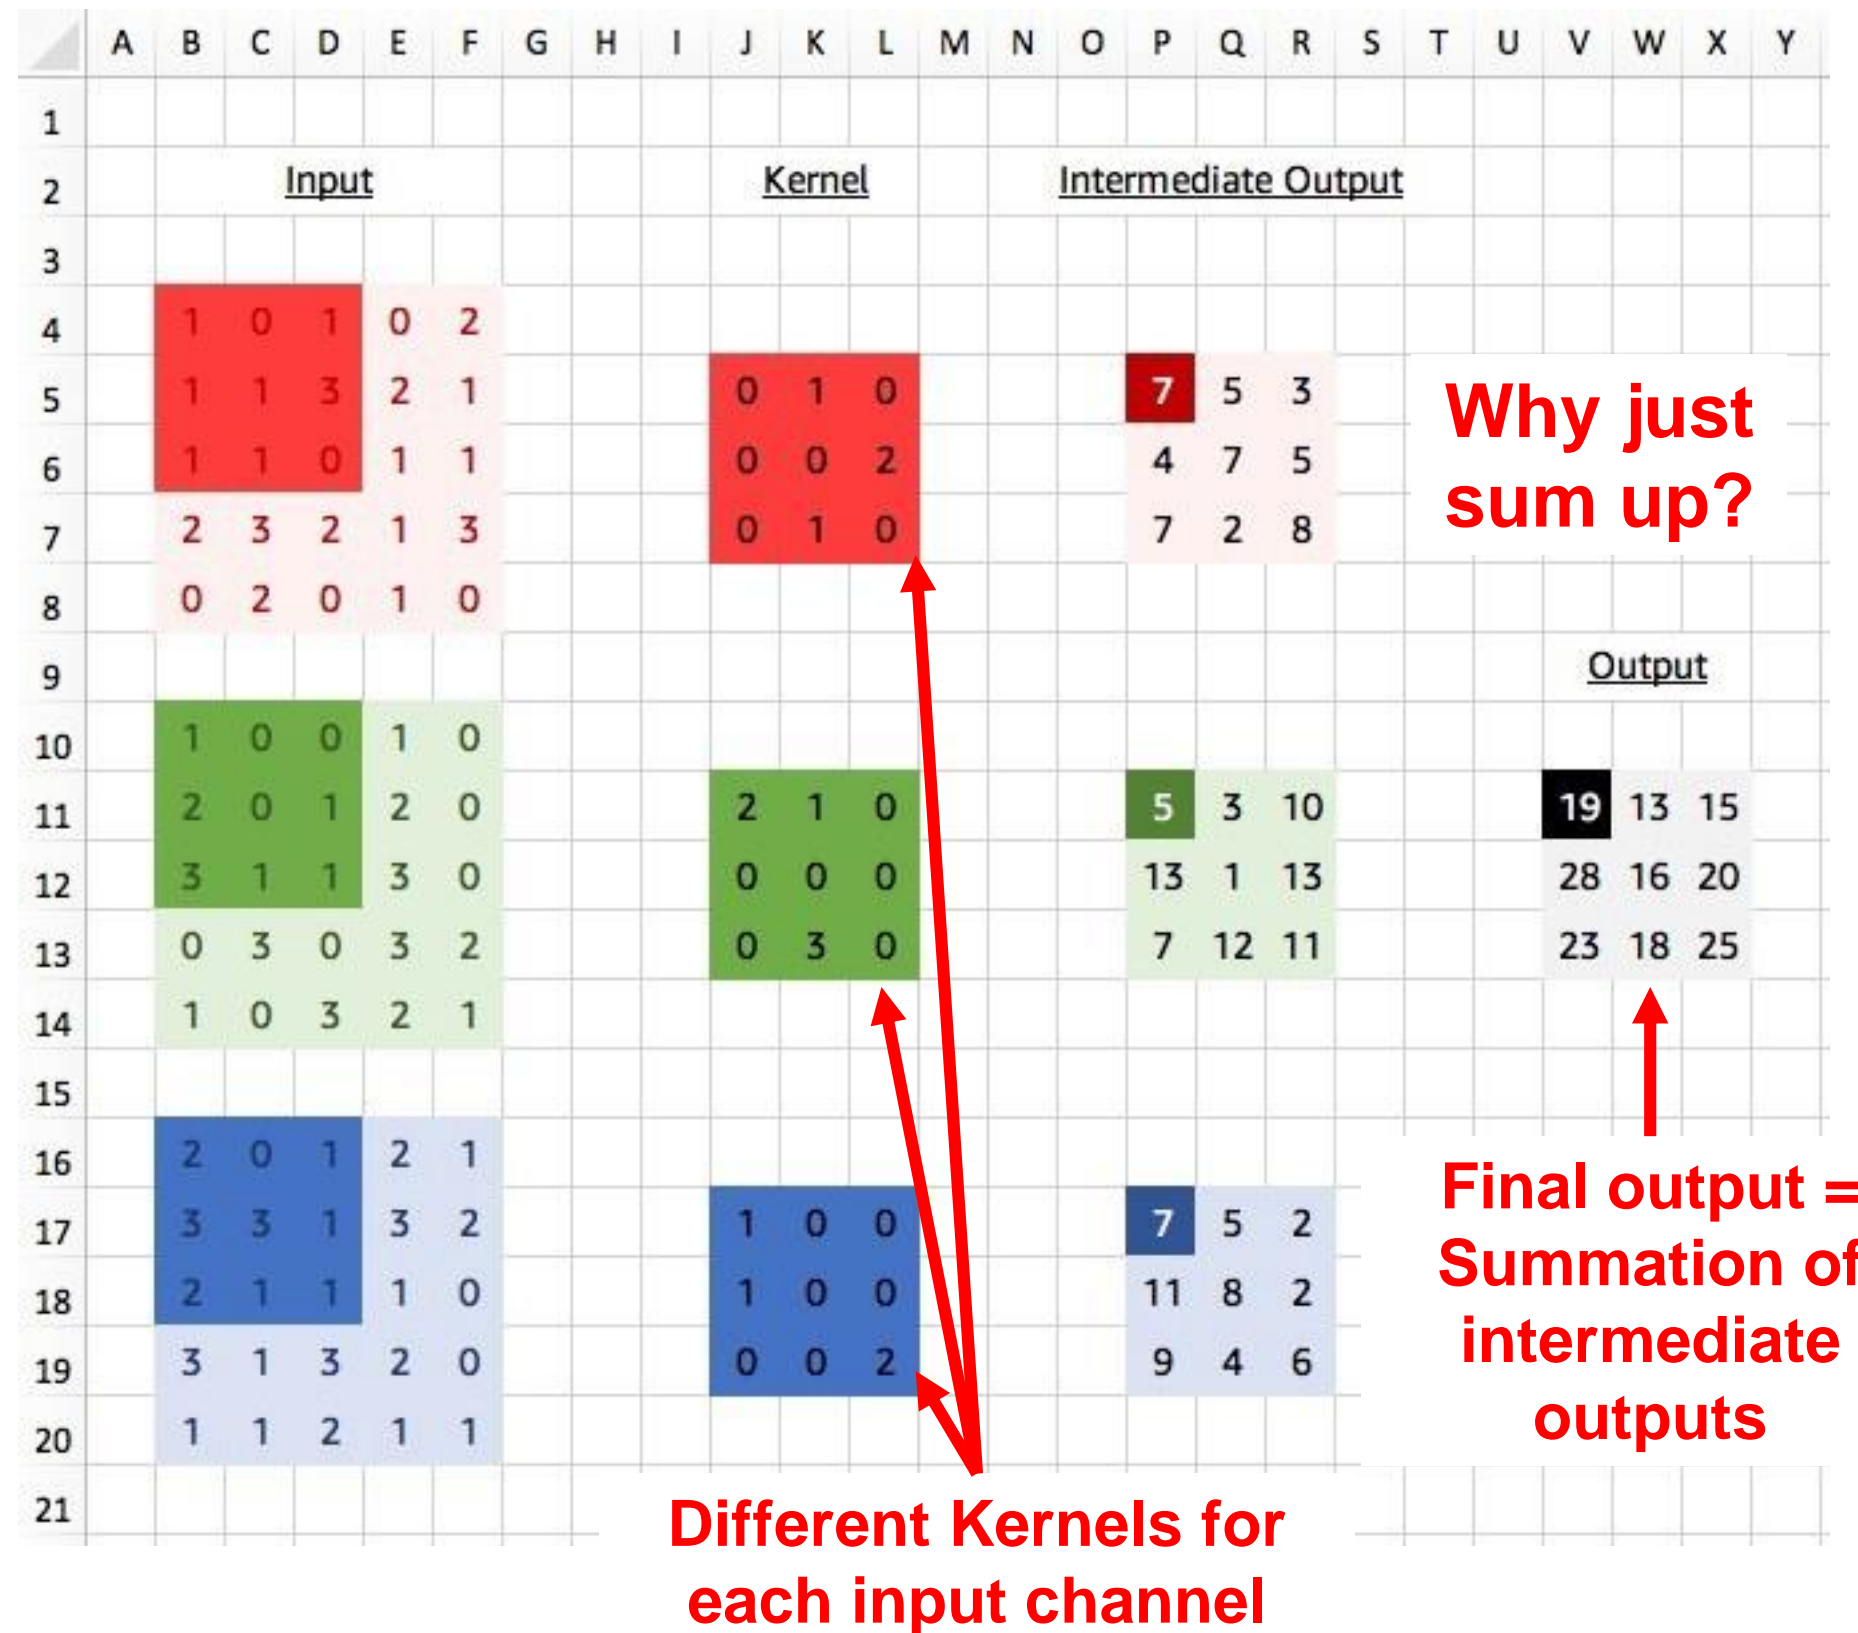
Source: <https://medium.com/apache-mxnet/convolutions-explained-with-ms-excel-465d6649831c>

# 2D convolution

Multi-channel

## Note:

- We do not decide the values of the Kernels
- The Kernel values are updated by the algorithms based on the training data
- We decide only the size of the Kernels



Source: <https://medium.com/apache-mxnet/convolutions-explained-with-ms-excel-465d6649831c>

# Max pooling

The original

## Convolutional vs max-pooling layers:

- **Convolutional layers extract features**
  - big Kernels extract obvious features,
  - small ones extract more detailed features
- **Max Pooling layers help to downscale the feature maps so that only features with significant weights are brought over to the next layer**

Single depth slice

|   |   |   |   |
|---|---|---|---|
| 1 | 1 | 2 | 4 |
| 5 | 6 | 7 | 8 |
| 3 | 2 | 1 | 0 |
| 1 | 2 | 3 | 4 |

max pool with 2x2 filters  
and stride 2



|   |   |
|---|---|
| 6 | 8 |
| 3 | 4 |

Source: <http://cs231n.github.io/convolutional-networks/>

# Max pooling

With situation

|   |   |   |   |   |
|---|---|---|---|---|
| 1 | 4 | 4 | 5 | 6 |
| 3 | 9 | 2 | 3 | 2 |
| 8 | 1 | 6 | 0 | 7 |
| 0 | 3 | 2 | 1 | 1 |

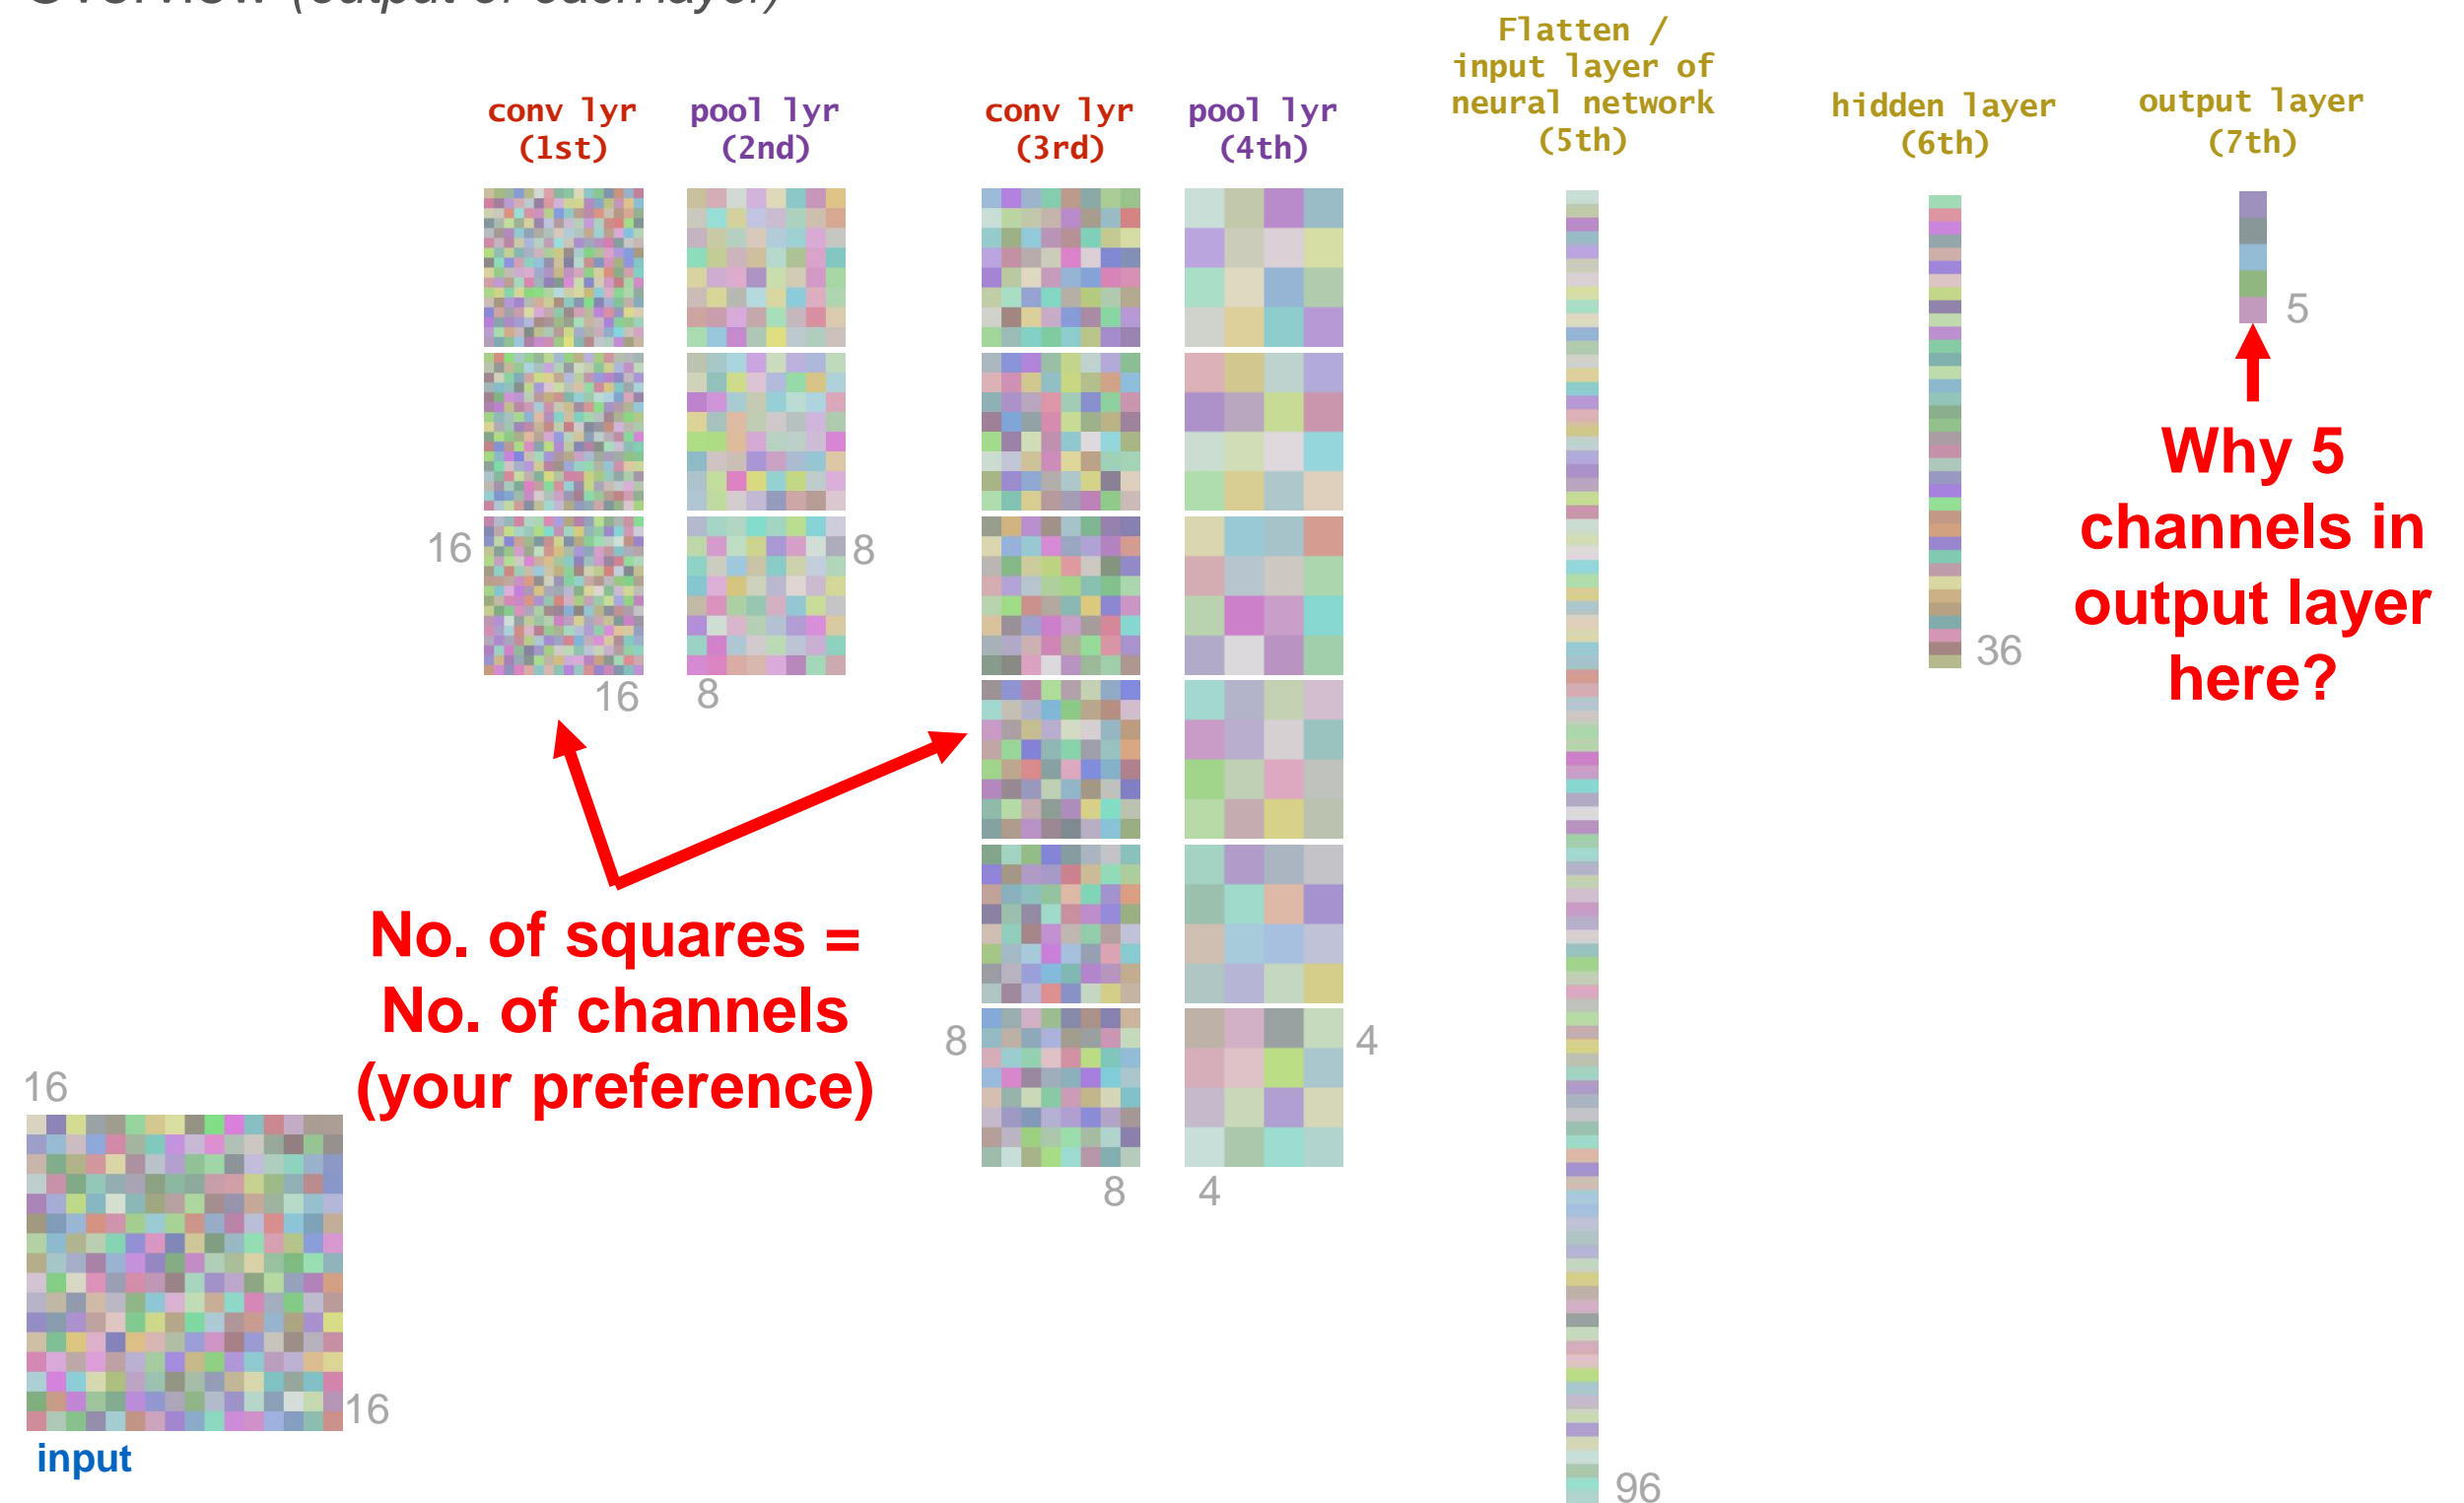


|   |   |
|---|---|
| 9 | 5 |
| 8 | 6 |

Source: <https://software.intel.com/en-us/daal-programming-guide-2d-max-pooling-forward-layer>

# Convolutional neural network

Overview (output of each layer)









# The making of ...

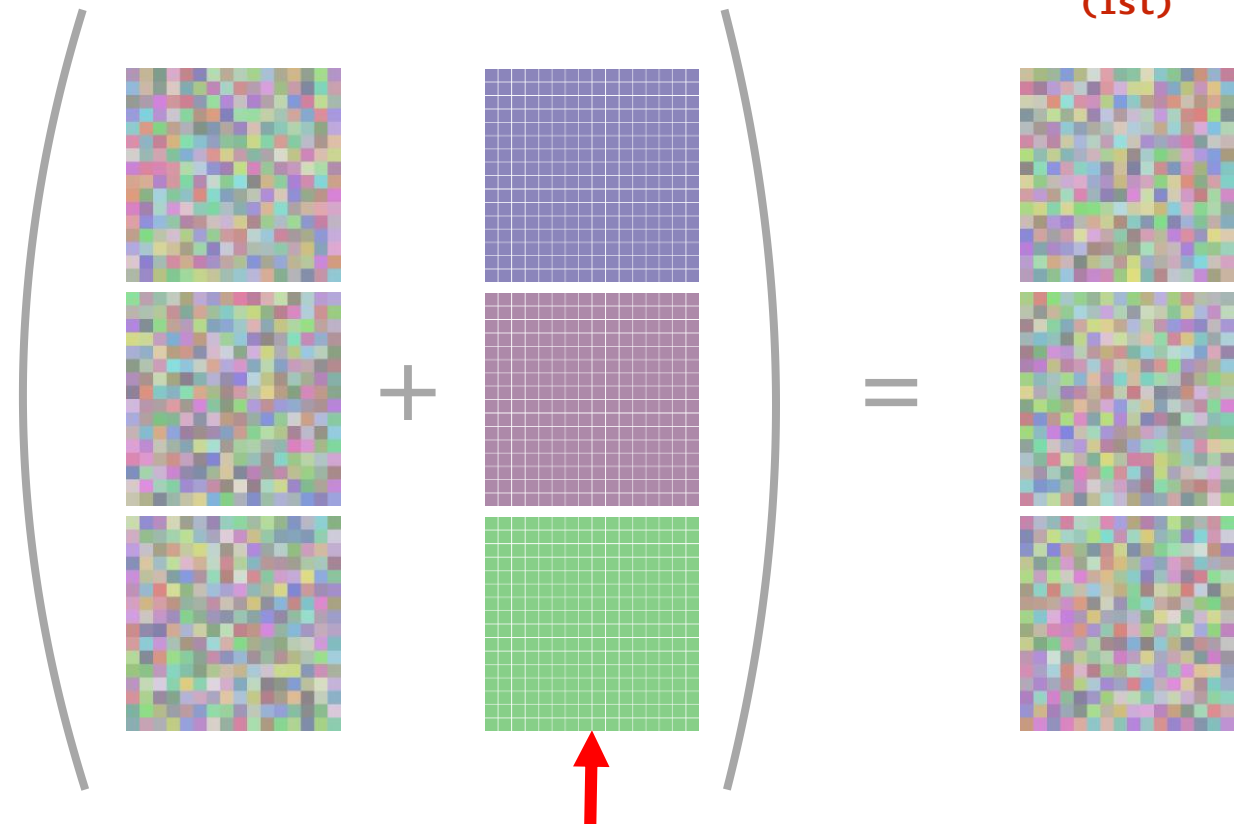
The first convolutional layer  
(part 2)

- Add bias to each convolution output, and apply activation function to get the final output for the convolutional layer

Can be ReLu or sigmoid function; the function is applied on each value in the matrix



activation


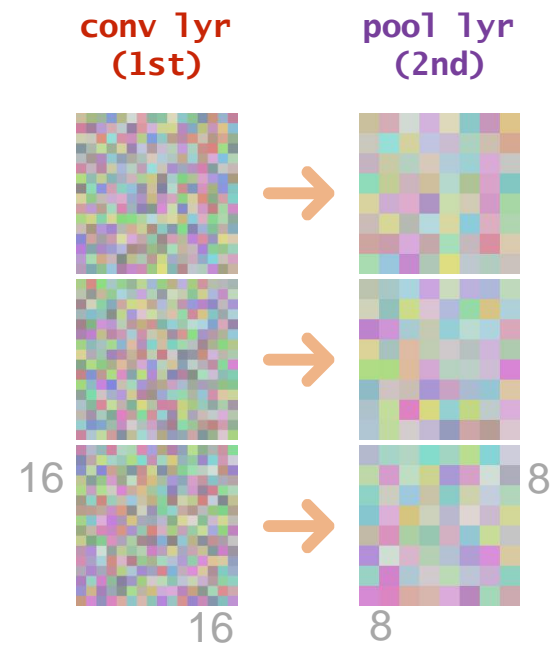


bias is a matrix and is the same throughout each matrix but different for various matrices

# The making of ....

The pooling layer

- Apply 2 x 2 max-pooling (stride 2) on the outputs from the first convolutional layer




# The making of ...


The second convolutional layer (part 1)

- Performs 6 separate multi-channel 2D convolutions (with padding) to generate 6 convolution outputs

pool 1yr  
(2nd)



1 set of 3 various Kernels  
for each input channel =  
3 immediate outputs are  
summed up to get one  
output channel



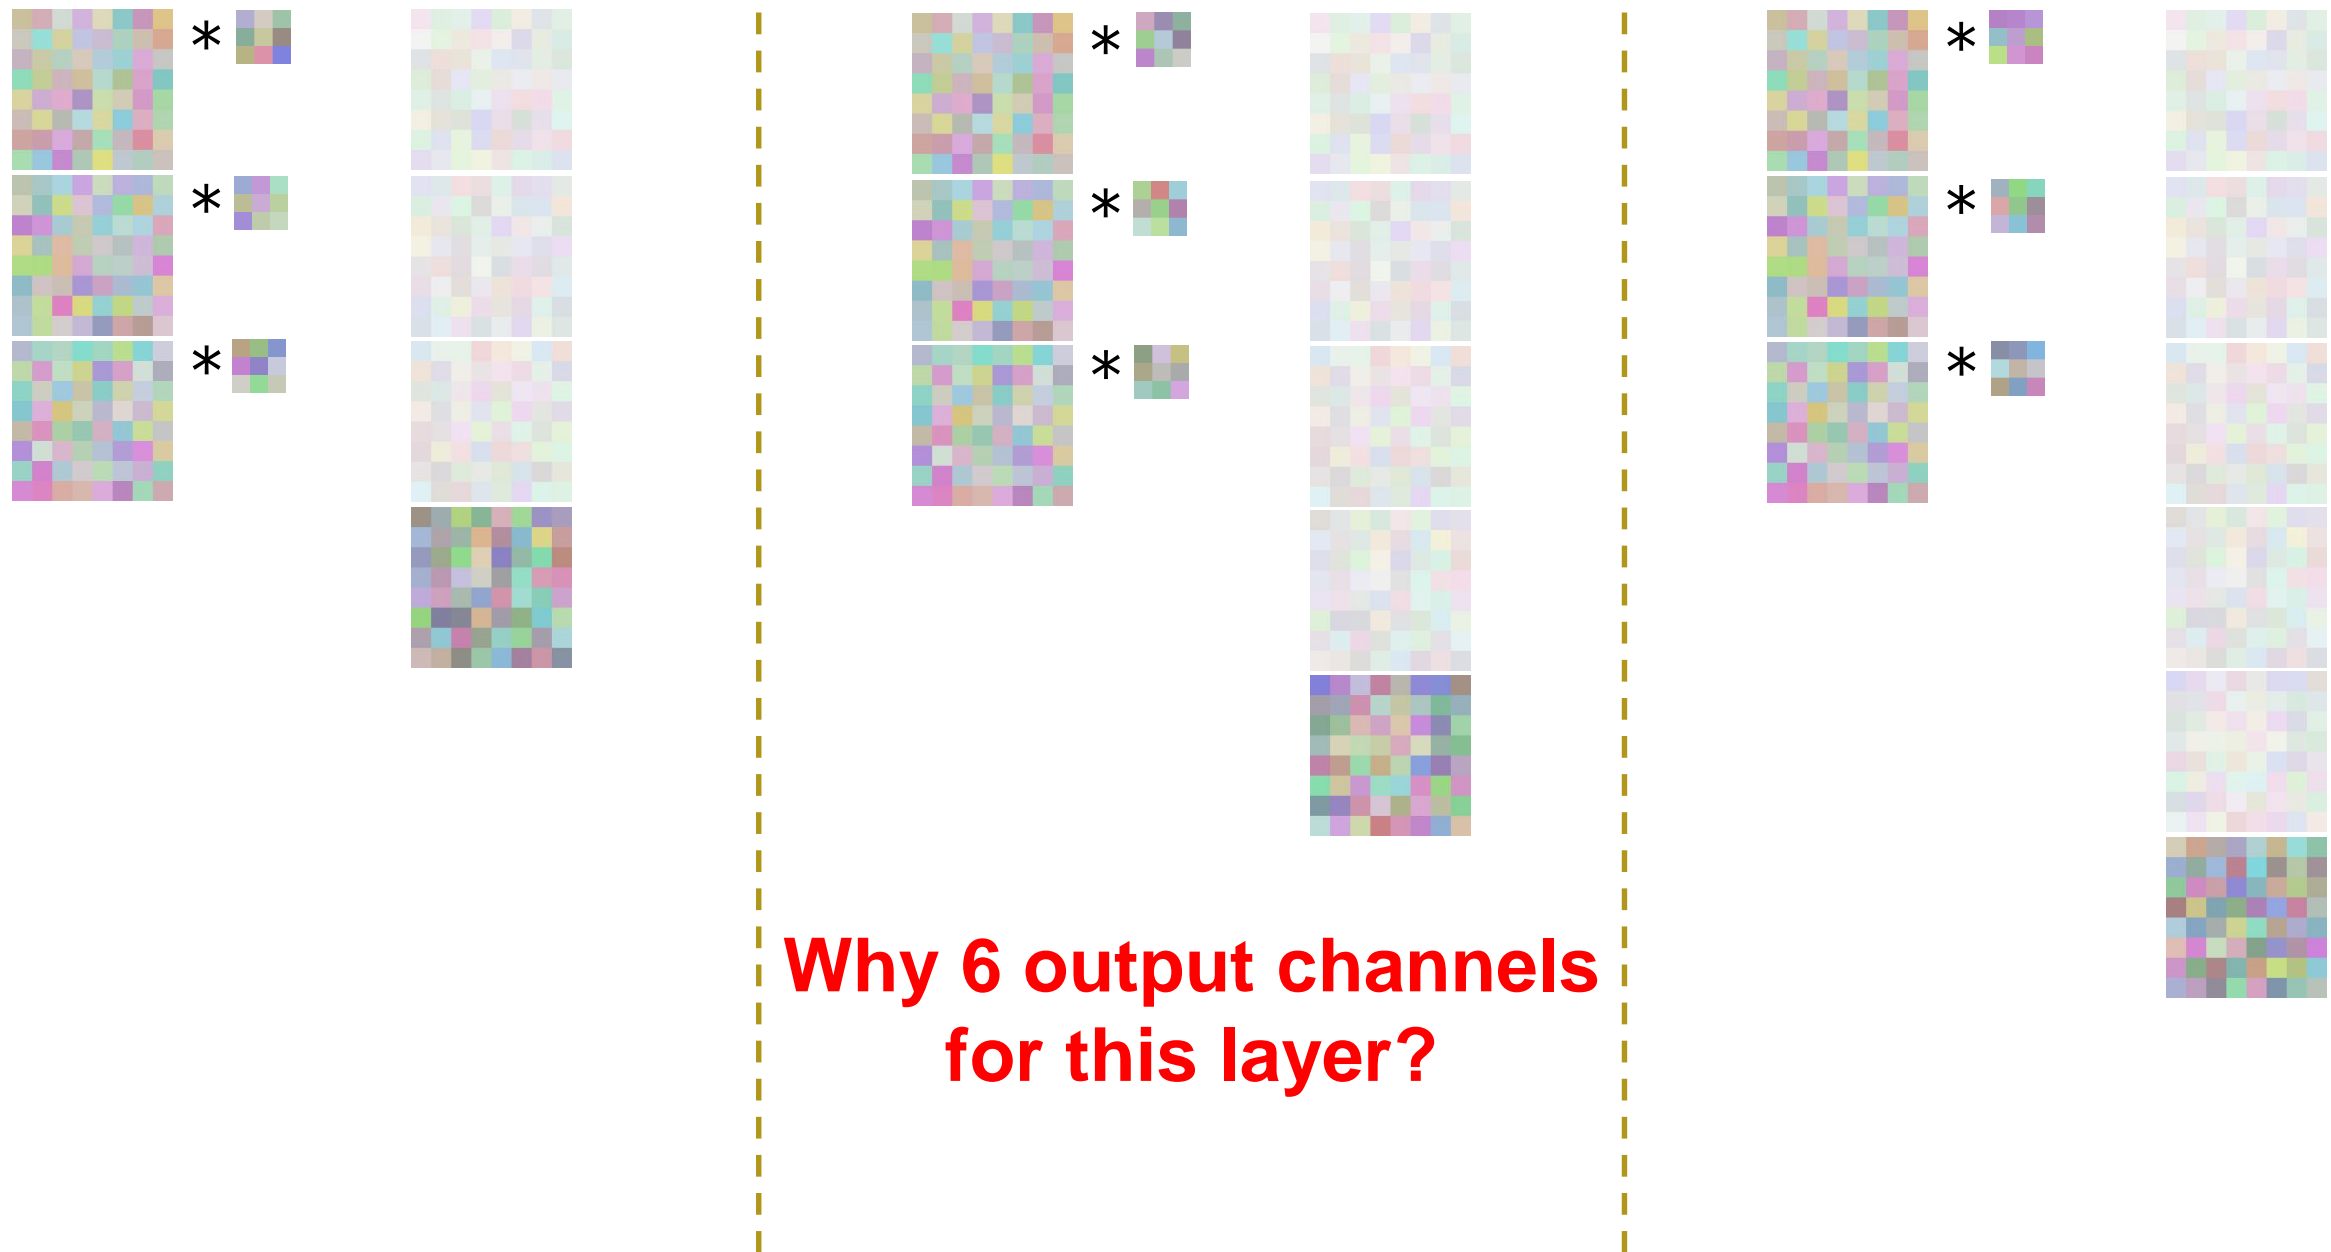
Do the same procedure to produce other  
output channels with different sets of  
Kernels

# The making of ...

The second convolutional layer (part 2)

- Performs 6 separate multi-channel 2D convolutions (with padding) to generate 6 convolution outputs

pool 1yr  
(2nd)




**Why 6 output channels  
for this layer?**

# The making of ...

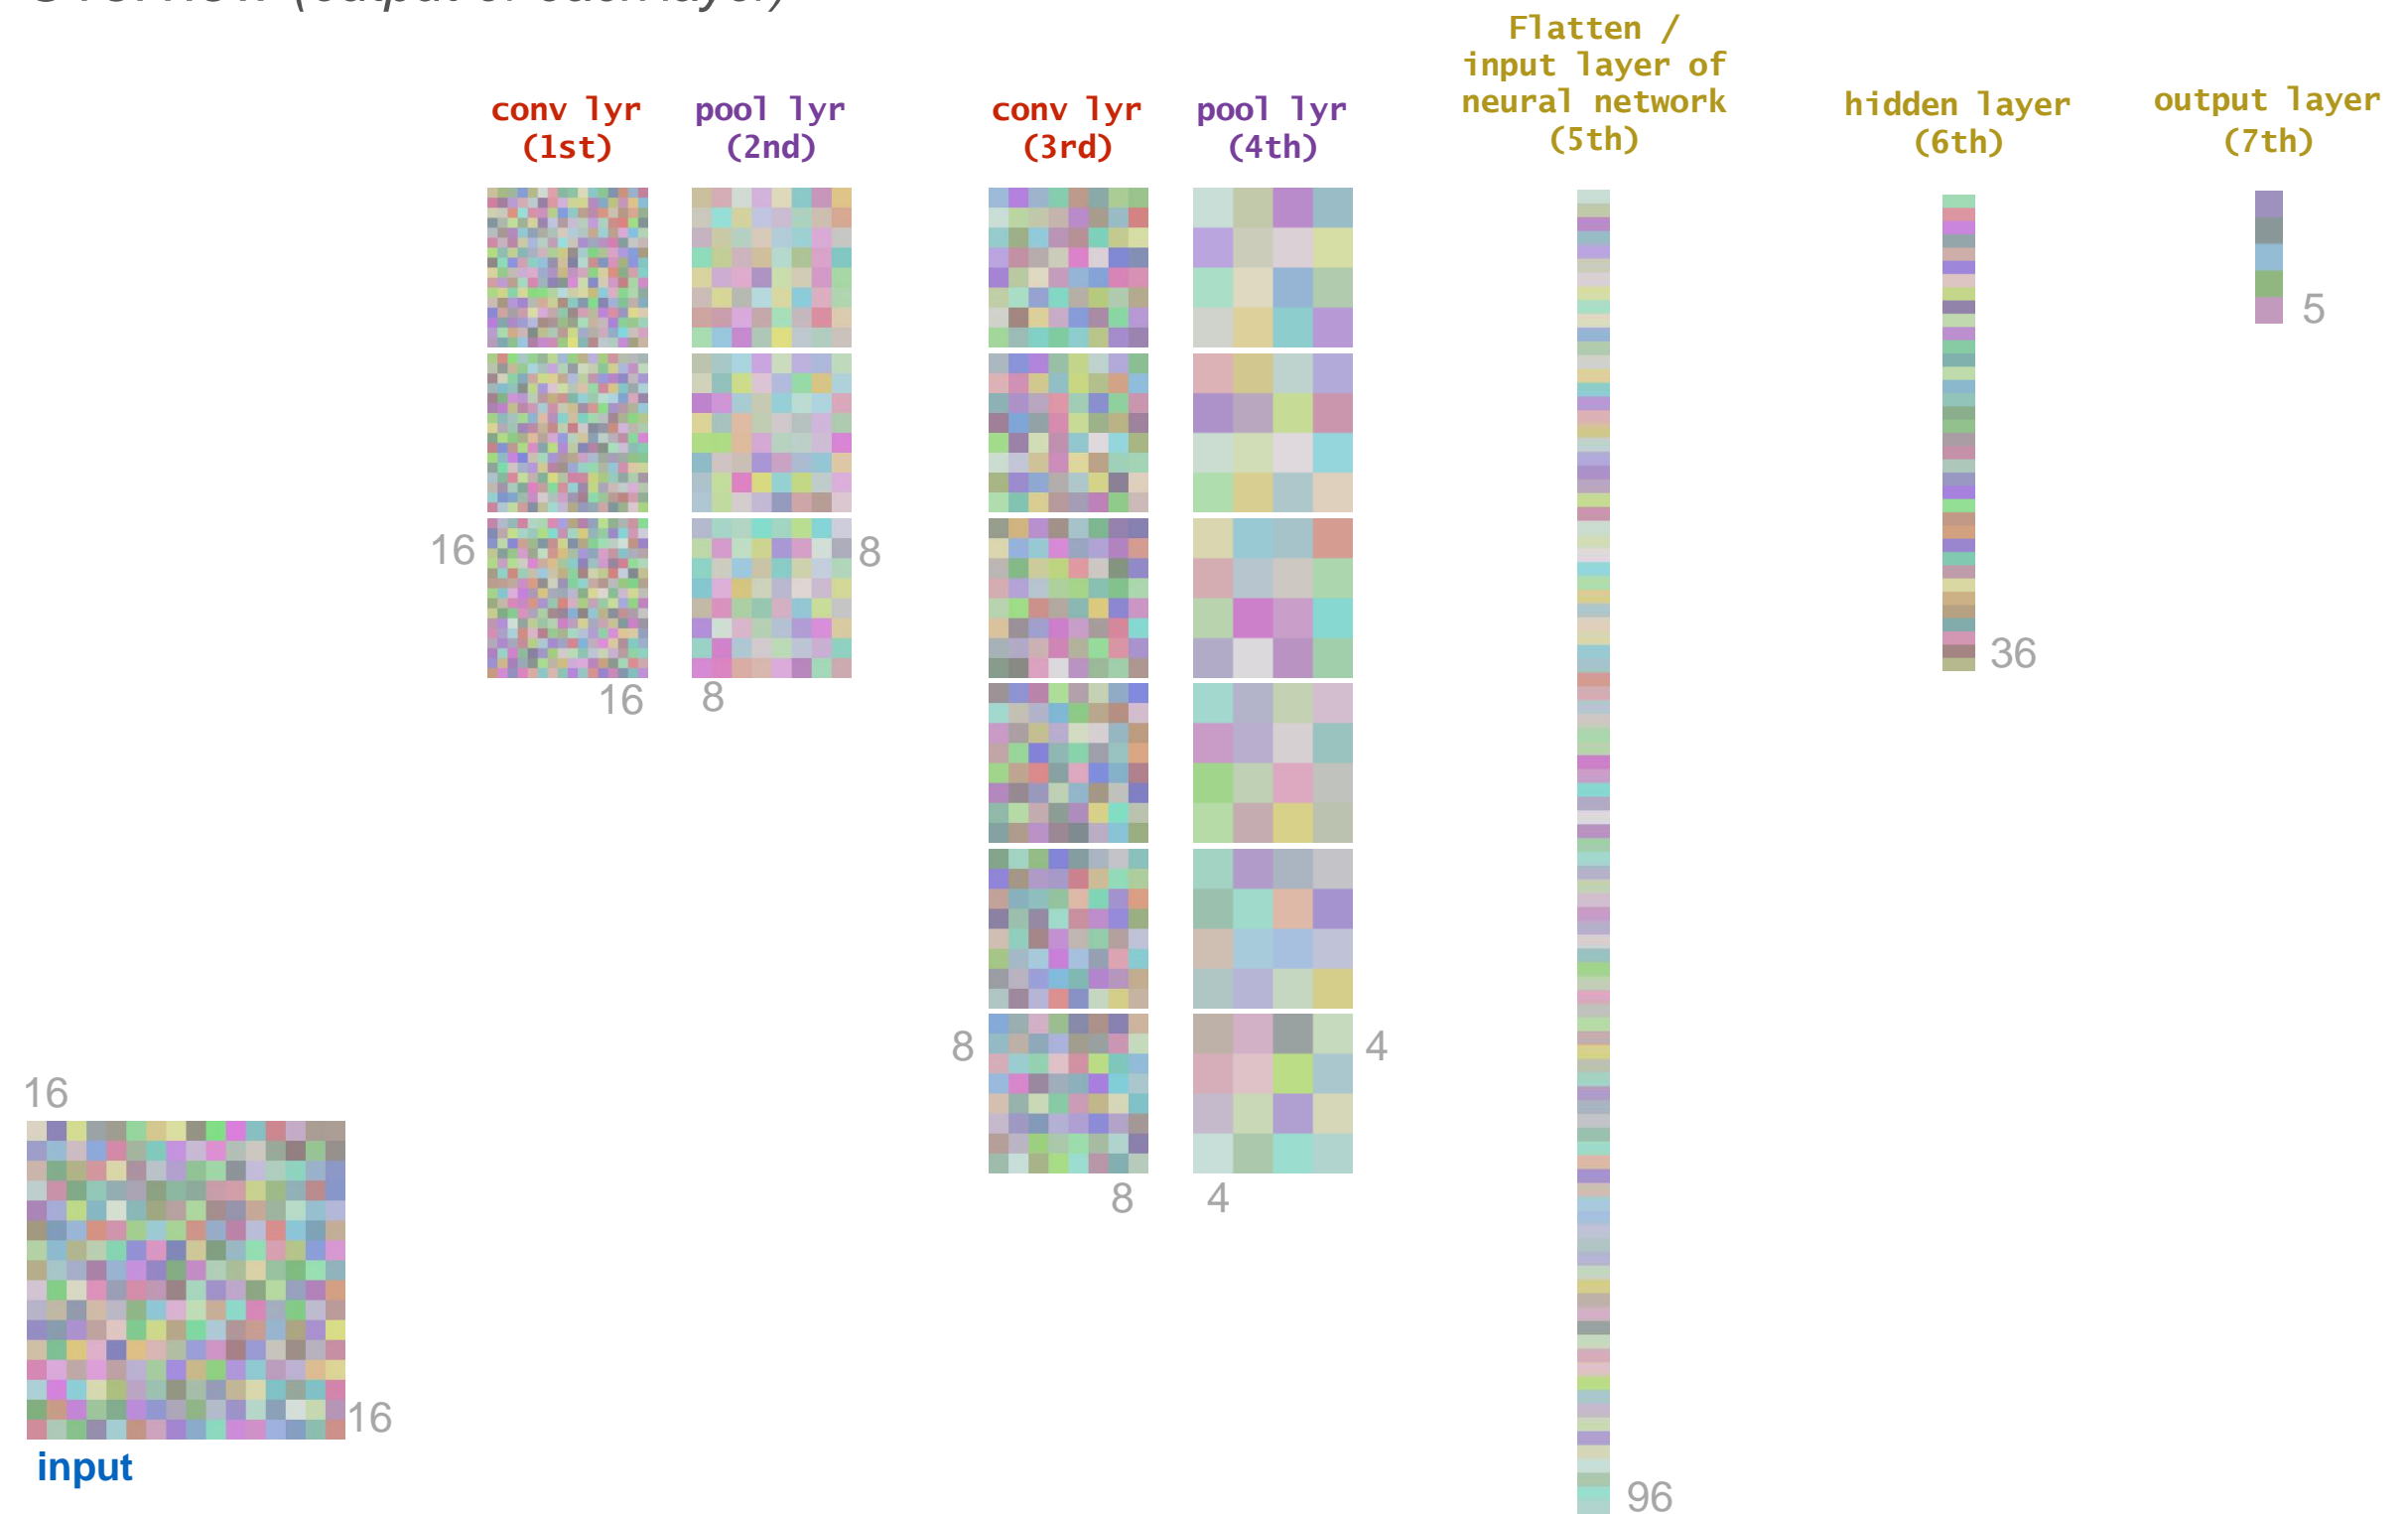
The first convolutional layer  
(part 3)

- Add bias to each intermediate output, and apply activation function to get the final output for the convolutional layer




# Convolutional neural network

Overview (*output of each layer*)



# Convolutional neural network

Dropout




**We will apply dropout in  
the workshop later on**

Source: <https://stats.stackexchange.com/questions/201569/difference-between-dropout-and-dropconnect>



# Convolutional neural network

## Dropout



- Very popular method to **regularize neural networks**; effective in preventing overfitting
- **Concept:**
  - Approximates training a large number of neural networks with different architectures in parallel
  - Every unit of the neural network (except output layer) is given the probability  $p$  of being temporarily ignored/muted (i.e. “dropped out”) in calculations
  - Hyper parameter  $p$  is called **dropout rate** and very often its default value is set to 0.5
  - In each iteration, the neurons are randomly selected according to the assigned probability. As a result, each time we work with a smaller neural network

Source: <https://towardsdatascience.com/preventing-deep-neural-network-from-overfitting-953458db800a>

**Any fans of Japan?**

**RECAP:**

**(3 basic components in deep learning learnt)**

**1. Convolutions**

**2. Pooling**

**3. Dropout**

# Cursive Kuzushiji

Automated solution?

**Workshop objective:**  
**To classify hand drawn Japanese characters**

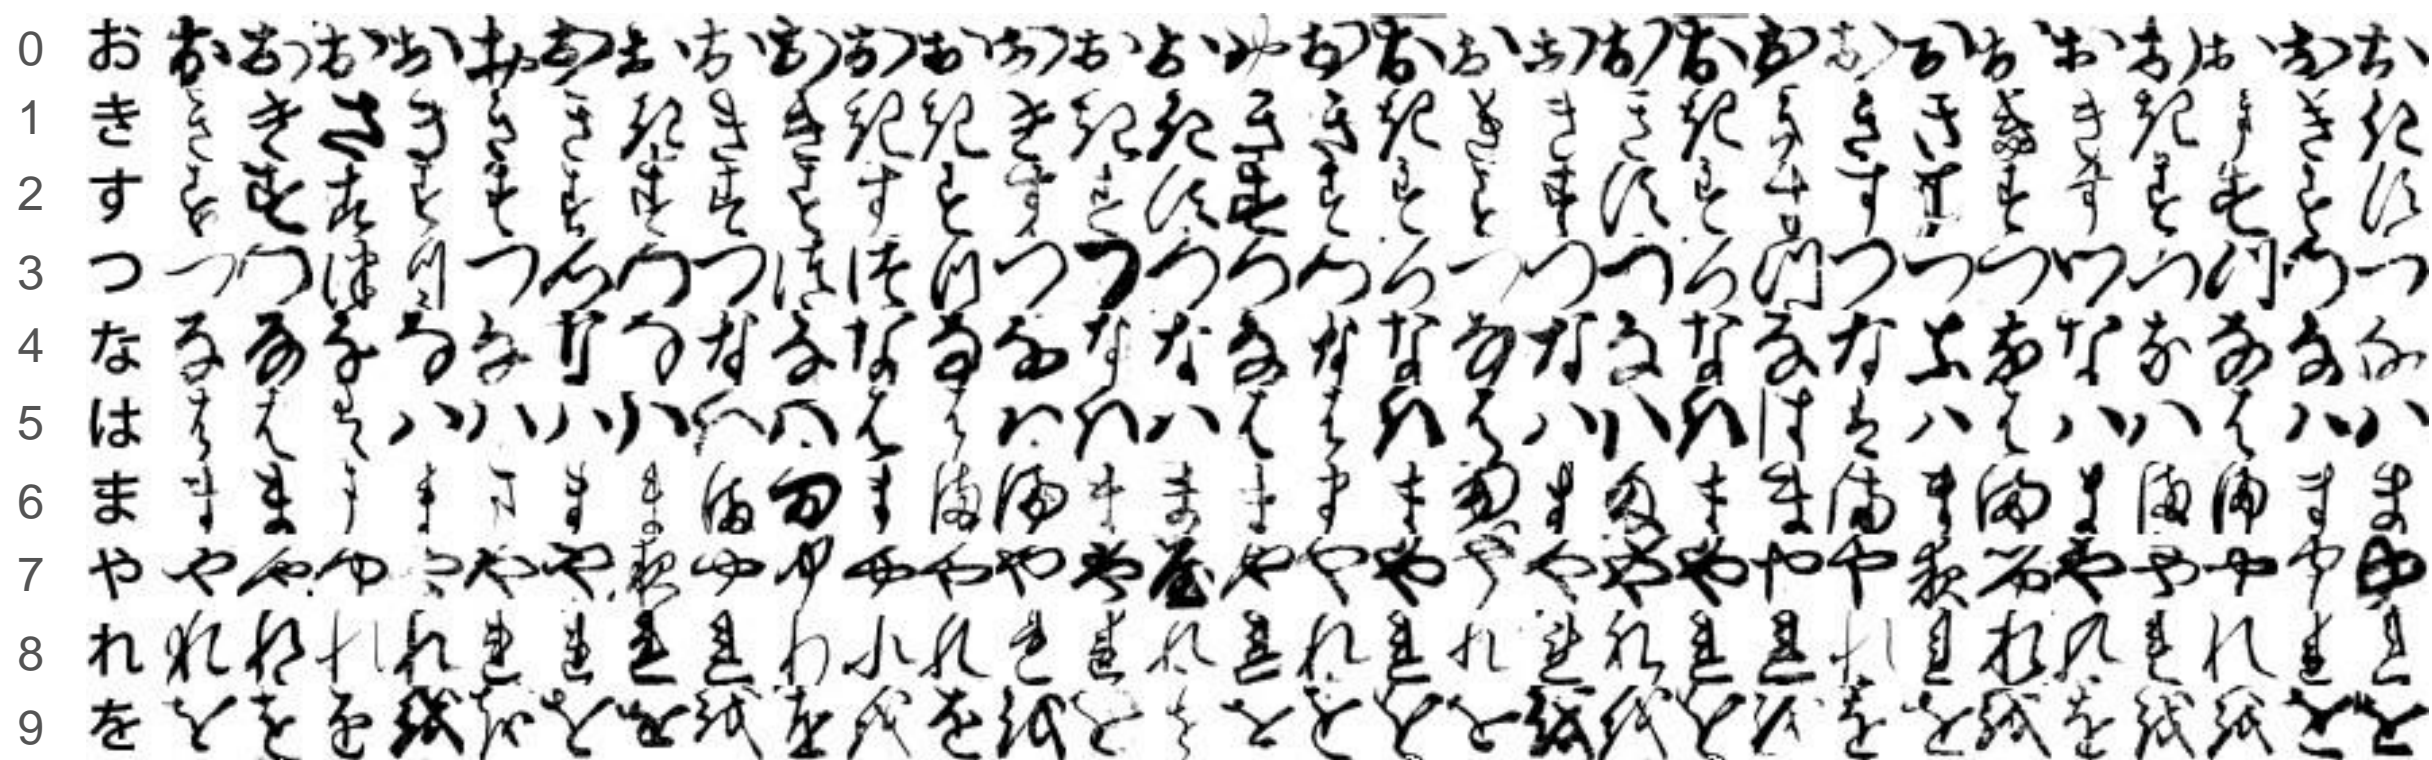


Source: <https://arxiv.org/pdf/1812.01718.pdf>



# Kuzushiji MNIST

Another 'MNIST' alternative



Source: [https://github.com/rois-codh/kmnist/blob/master/images/kmnist\\_examples.png](https://github.com/rois-codh/kmnist/blob/master/images/kmnist_examples.png)

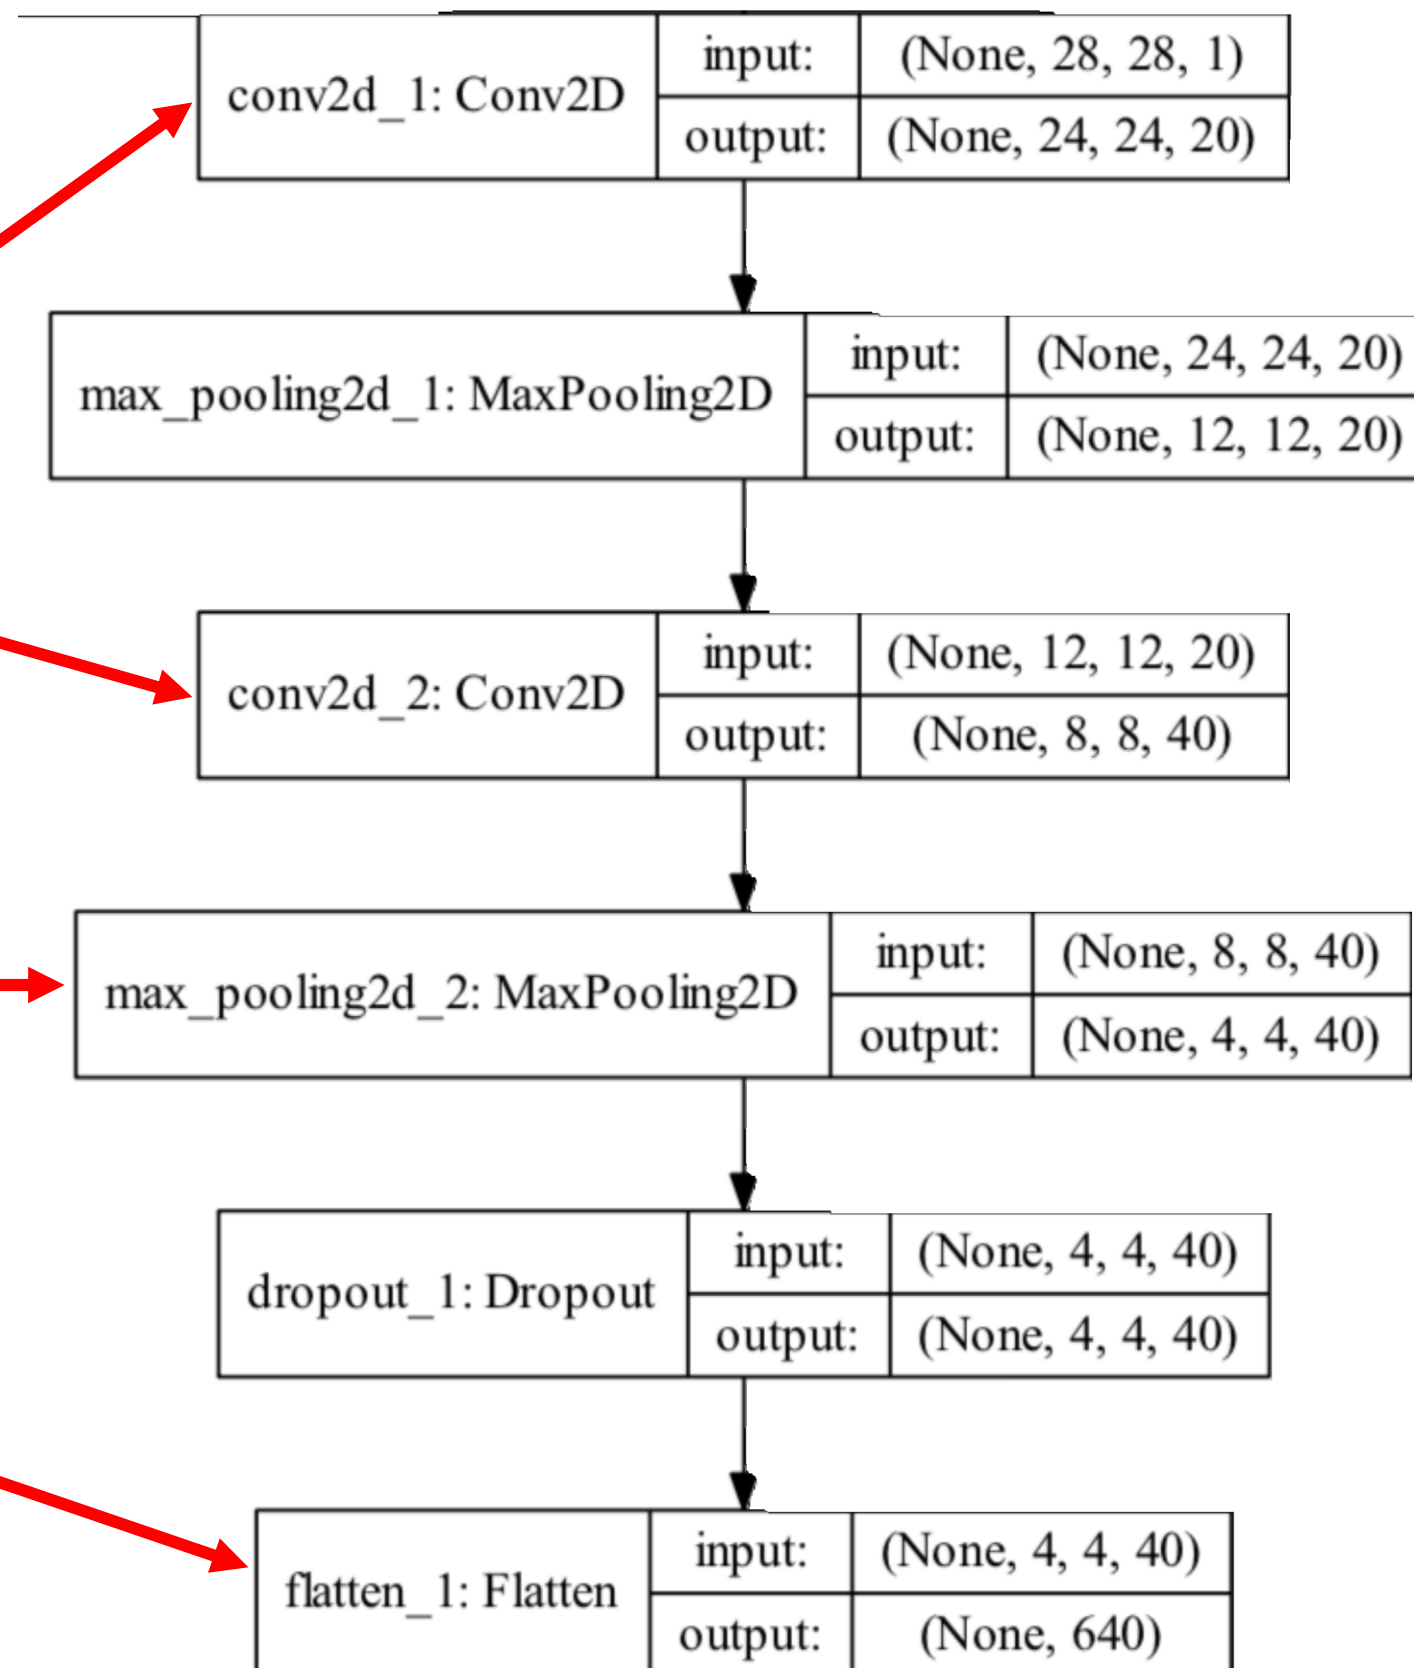
# Kuzushiji MNIST

The basic model, part 1

**Filter Kernel size of 5x5  
for all Conv Layers; to  
be shown later**

**Max pooling of size 2x2  
will reduce the matrix size  
by half; no. of channels  
does not change**

**When flatten, the layer  
becomes a 1D array;  
 $4 \times 4 \times 40 = 640$**

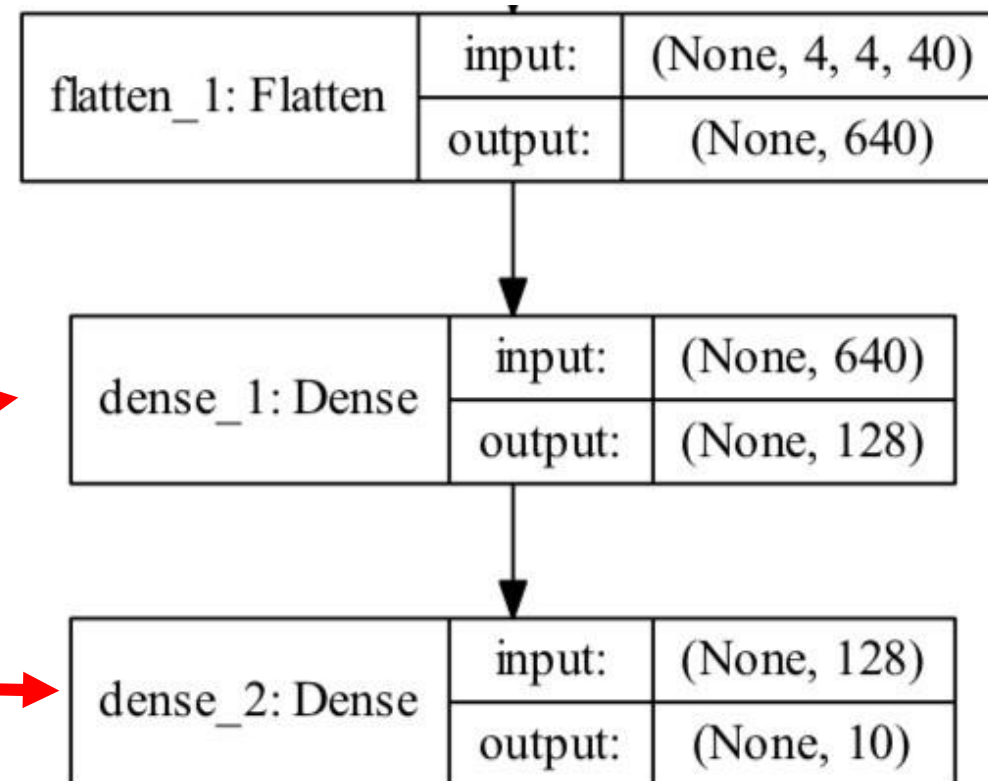


# Kuzushiji MNIST

The basic model, part 1

**Hidden Layer**

**Output Layer**



# Kuzushiji MNIST

The main layout for the code

1. Import libraries
2. Matplotlib setup
3. Data preparation
4. Define model
5. Train model
6. Test model



# Kuzushiji MNIST

## 1. Import libraries, part 1

- numpy for matrix manipulation;
- sklearn for measuring performance;
- matplotlib to show image and plot result;
- os for path manipulation

```
> import numpy as np
> import sklearn.metrics as metrics
> import matplotlib.pyplot as plt
> import os
```

# Kuzushiji MNIST

## 1. Import libraries, part 2

- Import all the Keras functions that we are going to use in this problem; note that we are using Keras function under the tensorflow; not using the keras directly

```
> from tensorflow.keras.callbacks import ModelCheckpoint, CSVLogger  
> from tensorflow.keras.models import Sequential  
> from tensorflow.keras.layers import Dense  
> from tensorflow.keras.layers import Dropout  
> from tensorflow.keras.layers import Flatten  
> from tensorflow.keras.layers import Conv2D  
> from tensorflow.keras.layers import MaxPooling2D  
> from tensorflow.keras.utils import to_categorical
```

**To save the model**

**To train and test data using the model**

**Approach to build our model for this WS; other approaches will be covered in PRMLS course**

**Function that allows us to convert our labels from integer into a one hot encoding type**

# Kuzushiji MNIST

## 2. Matplotlib setup, part 1

- First three lines setup the font manager, so that we can display Japanese words correctly in later usage
- Use 'ggplot' style to plot our training and testing result

```
> from matplotlib import font_manager as fm
> fpath          = os.path.join(os.getcwd(), "ipam.ttf")
> prop           = fm.FontProperties(fname=fpath)

> plt.style.use('ggplot')
> plt.rcParams['ytick.right']      = True
> plt.rcParams['ytick.labelright']= True
> plt.rcParams['ytick.left']      = False
> plt.rcParams['ytick.labelleft'] = False
> plt.rcParams['font.family']     = 'Arial'
```

# Kuzushiji MNIST

## 2. Matplotlib setup, part 2

- Create a function that can display gray scale image correctly

```
> def grayplt(img, title=''):
    plt.axis('off')
    if np.size(img.shape) == 3:
        plt.imshow(img[:, :, 0], cmap='gray', vmin=0, vmax=1)
    else:
        plt.imshow(img, cmap='gray', vmin=0, vmax=1)
    plt.title(title, fontproperties=prop)
    plt.show()
```

**The key is to set vmin and vmax 0 and 1 in order to display the gray scale correctly**

**Otherwise imshow will rescale our gray scale images, which is not desirable for this application**

# Kuzushiji MNIST

## 3. Data preparation, part 1

- Load train and test data; load train and test labels
- Rescale data to float, range from 0 to 1

**Use numpy function to load the data (in npz format in numpy)**

```
> trDat = np.load('kmnist-train-imgs.npz')['arr_0']
> trLbl = np.load('kmnist-train-labels.npz')['arr_0']
> tsDat = np.load('kmnist-test-imgs.npz')['arr_0']
> tsLbl = np.load('kmnist-test-labels.npz')['arr_0']
```

**Provided data is in uint8 format (unsigned 8-bit integer; range 0~255; this is too large!)**

```
> trDat = trDat.astype('float32')/255
> tsDat = tsDat.astype('float32')/255
```

**Hence, convert data type to float for both tr and ts data (rescale range to 0~1; optimal range)**

```
> imgrows = trDat.shape[1]
> imgcols = trDat.shape[2]
```

# Kuzushiji MNIST

## 3. Data preparation, part 2

**60000 images for tr;  
rows & col is 28x28**

**10000 images for ts;  
rows & col is 28x28**

- The current shape for trDat is (60000, 28, 28)
- The current shape for tsDat is (10000, 28, 28)
- Need to be reshaped into the form of (samples, width, height, channel) to fit into Keras API

```
> trDat      = trDat.reshape(trDat.shape[0],  
                             imgrows,  
                             imgcols,  
                             1)  
> tsDat      = tsDat.reshape(tsDat.shape[0],  
                             imgrows,  
                             imgcols,  
                             1)
```

**Reshaping functions to  
include the channel;  
necessary step to meet  
the requirements for the  
Keras API**

# Kuzushiji MNIST


## 3. Data preparation, part 3

- One-hot encode the train and test label information; get the number of classes in the labels

Use the `to_categorical` function to perform the one-hot encoding

```
trLb1 = to_categorical(trLb1)
tsLb1 = to_categorical(tsLb1)
> num_classes = tsLb1.shape[1]
```

Preparation done! Now we want to define the model



Example of a 4 class label that are converted into a one-hot encoding type

Source: <https://arxiv.org/pdf/1812.01718.pdf>



# Kuzushiji MNIST

## 4. Define model, part 1

**Main Thing to do for today's workshop!**

```
> seed = 29
> np.random.seed(seed)
```

```
> modelname = 'wks5_1a'
```

```
> def createModel():
```

```
    model = Sequential()
```

```
    model.add(Conv2D(20, (5, 5), input_shape=(28, 28, 1), activation='relu'))
```

```
    model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
    model.add(Conv2D(40, (5, 5), activation='relu'))
```

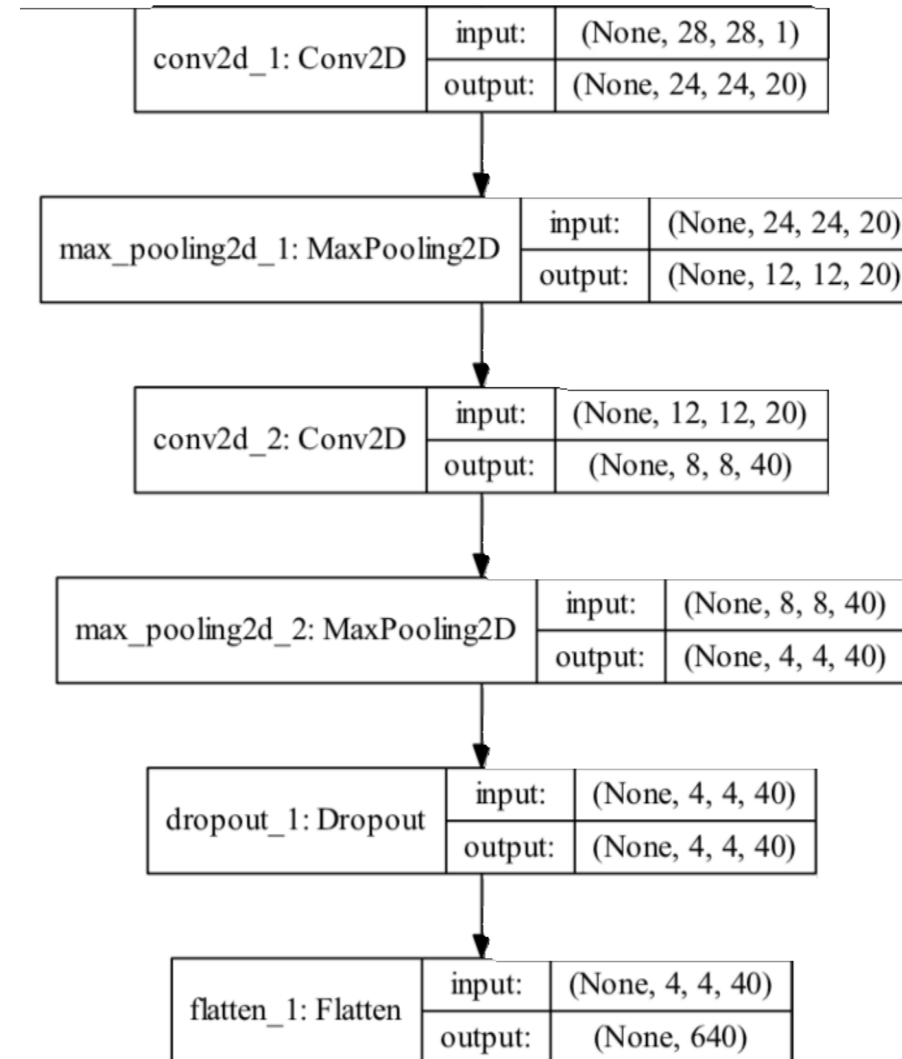
```
    model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
    model.add(Dropout(0.2))
```

```
    model.add(Flatten())
```

```
    model.add(Dense(128, activation='relu'))
```

```
    ....
```



**Note that input shape has to be defined only for the 1<sup>st</sup> Conv2D layer**

**Dropout rate = 0.2**

# FAQs

1. Where to put the Dropouts?
2. Max Pooling size strictly 2x2?
3. Bigger vs Smaller Kernel size?
4. How many channels/neurons should I put for each CNN/Dense layer?
5. How many CNN/Dense layers should I put?
6. How to enable GPU on colab?
7. How to put padding? (to be answered next few slides)
8. How do I change the stride value? (to be answered next few slides)


# Kuzushiji MNIST

## 4. Define model, part 1

```
> seed = 29
> np.random.seed(seed)

> modelName = 'wks5_1a'
> def createModel():
    model = Sequential()
    model.add(Conv2D(20, (5, 5), input_shape=(28, 28, 1), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Conv2D(40, (5, 5), activation='relu', padding='same/valid'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.2))
    model.add(Flatten())
    model.add(Dense(128, activation='relu'))
    .....
```

## Changing Padding Type



# Kuzushiji MNIST

## 4. Define model, part 1

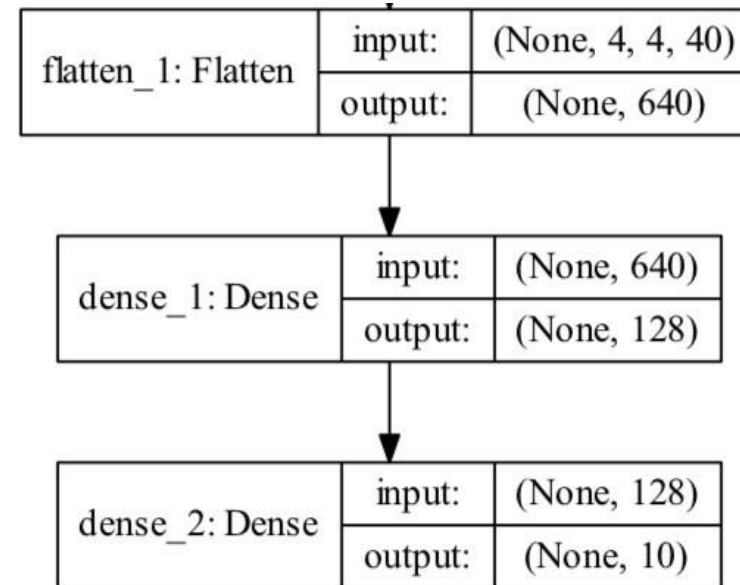
## Changing Stride Value

```
> seed = 29
> np.random.seed(seed)

> modelname = 'wks5_1a'
> def createModel():
    model = Sequential()
    model.add(Conv2D(20, (5, 5), input_shape=(28, 28, 1), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
    model.add(Conv2D(40, (5, 5), strides=(2, 2), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.2))
    model.add(Flatten())
    model.add(Dense(128, activation='relu'))
    ....
```

# Kuzushiji MNIST

## 4. Define model, part 1




```
> seed = 29
> np.random.seed(seed)

> modelname = 'wks5_1a'
> def createModel():
    model = Sequential()
    model.add(Conv2D(20, (5, 5), input_shape=(28, 28, 1), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Conv2D(40, (5, 5), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.2))
    model.add(Flatten())
    model.add(Dense(128, activation='relu'))
    model.add(Dense(num_classes, activation='softmax'))
    model.compile(loss='categorical_crossentropy', optimizer='adam',
                  metrics=['accuracy'])
    return model
```

**Dense layers perform the main classification tasks**

# Kuzushiji MNIST

## 4. Define model, part 1



```
> seed = 29
> np.random.seed(seed)
```

```
> modelname = 'wks5_1a'
```

```
> def createModel():
```

```
    model = Sequential()
```

```
    model.add(Conv2D(20, (5, 5), input_shape=(28, 28, 1), activation='relu'))
```

```
    model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
    model.add(Conv2D(40, (5, 5), activation='relu'))
```

```
    model.add(MaxPooling2D(pool_size=(2, 2)))
```

```
    model.add(Dropout(0.2))
```

```
    model.add(Flatten())
```

```
    model.add(Dense(128, activation='relu'))
```

```
    model.add(Dense(num_classes, activation='softmax'))
```

```
    model.compile(loss='categorical_crossentropy', optimizer='adam',
```

```
                  metrics=['accuracy'])
```

```
    return model
```

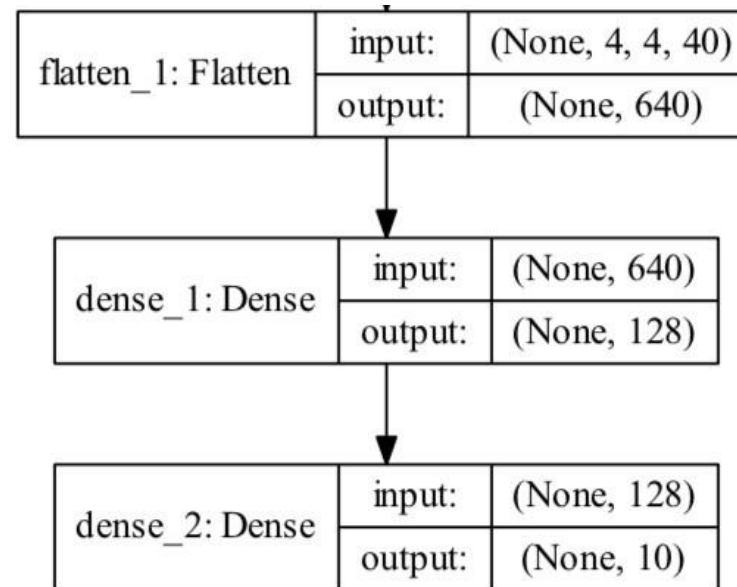
**Softmax enables you to identify which output is true within the output layer (in terms of probabilities)**

# Kuzushiji MNIST

## 4. Define model, part 1

```
> seed = 29
> np.random.seed(seed)

> modelname = 'wks5_1a'
> def createModel():
    model = Sequential()
    model.add(Conv2D(20, (5, 5), input_shape=(28, 28, 1), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Conv2D(40, (5, 5), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))
    model.add(Dropout(0.2))
    model.add(Flatten())
    model.add(Dense(128, activation='relu'))
    model.add(Dense(num_classes, activation='softmax'))
    model.compile(loss='categorical_crossentropy', optimizer='adam',
                  metrics=['accuracy'])
    return model
```



**Optimizer = an algorithm that tells the framework how to update the weights and bias. Backbone of this is backpropagation; different optimizers have different ways to update the weights and bias**



# Kuzushiji MNIST

## 4. Define model, part 2

- 'model' for training; 'modelGo' for final evaluation

**Summary allows you monitor your model to check if your model construction is done correctly**

```
> model = createModel()  
> modelGo = createModel()  
  
model.summary()
```

Output example

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d_1 (Conv2D)              | (None, 24, 24, 20) | 520     |
| max_pooling2d_1 (MaxPooling2D) | (None, 12, 12, 20) | 0       |
| conv2d_2 (Conv2D)              | (None, 8, 8, 40)   | 20040   |
| max_pooling2d_2 (MaxPooling2D) | (None, 4, 4, 40)   | 0       |
| dropout_1 (Dropout)            | (None, 4, 4, 40)   | 0       |
| flatten_1 (Flatten)            | (None, 640)        | 0       |
| dense_1 (Dense)                | (None, 128)        | 82048   |
| dense_2 (Dense)                | (None, 10)         | 1290    |
| Total params: 103,898          |                    |         |
| Trainable params: 103,898      |                    |         |
| Non-trainable params: 0        |                    |         |

**Creating two models here (one for training and another for final evaluation); not necessary to do this but it is a good practice**

# Kuzushiji MNIST

## 4. Define model, part 3

- Create checkpoints to save model during training and save training data into csv

**Before training, we need to specify where/how we are saving the model**

```
> filepath = modelname + ".hdf5"
> checkpoint = ModelCheckpoint(filepath,
                               monitor='val_acc',
                               verbose=0,
                               save_best_only=True,
                               mode='max')
```

**Use ModelCheckpoint to save the model in the middle of the training process**

**Monitor model set to "validation accuracy"**

**Save when the validation accuracy is the max (i.e. mode = max)**

```
> csv_logger = CSVLogger(modelname + '.csv')
> callbacks_list = [checkpoint, csv_logger]
```

**Lastly, put the 2 objects (i.e. checkpoint and csv\_logger) into a list, named as callbacks\_list; these 2 objects (i.e. callbacks) will be called after each training epoch**

**Log the training and testing information into a CSV file via CSVlogger**

# Kuzushiji MNIST

## 5. Train model

- Training is only a single line

**Use the `model.fit` function to do the training**

```
model.fit(trDat,
          trLbl,
          validation_data=(tsDat, tsLbl),
          epochs=60,
          batch_size=128,
          callbacks=callbacks_list)
```

**Epochs is a hyperparameter that defines the number of times that the learning algorithm will work through the entire training dataset**

**Batch size is a hyperparameter that controls each time how many samples are taken at one go to train and update the weights**

Train on 60000 samples, validate on 10000 samples

Epoch 1/60

60000/60000 [=====] - 68s 1ms/sample - loss: 0.4707 - acc: 0.8539 - val\_loss: 0.4163 - val\_acc: 0.8689

Epoch 2/60

60000/60000 [=====] - 66s 1ms/sample - loss: 0.1603 - acc: 0.9509 - val\_loss: 0.3003 - val\_acc: 0.9125

Epoch 3/60

60000/60000 [=====] - 66s 1ms/sample - loss: 0.1068 - acc: 0.9673 - val\_loss: 0.2459 - val\_acc: 0.9290

Epoch 4/60

60000/60000 [=====] - 65s 1ms/sample - loss: 0.0798 - acc: 0.9751 - val\_loss: 0.2348 - val\_acc: 0.9352

Epoch 5/60

60000/60000 [=====] - 65s 1ms/sample - loss: 0.0653 - acc: 0.9794 - val\_loss: 0.2254 - val\_acc: 0.9406

.....

# Kuzushiji MNIST

## 6. Test model, part 1

- Use a new object to load the weights, and check the best accuracy

**After training, we need to do a final evaluation (no training) using a fresh model (i.e. modelGo); must be fresh to test the model with the trained weights**

```
> modelGo.load_weights(filepath)
> modelGo.compile(loss='categorical_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])
```

**Weights loaded from the trained model and the fresh model is recompiled**

# Kuzushiji MNIST


## 6. Test model, part 2

- Test the model, calculate the accuracy and confusion matrix

Use the `modelGo.predict` function to test the model using the testing data set

```
> predicts = modelGo.predict(tsDat)
> predout = np.argmax(predicts, axis=1)
```

Use this function to convert “predicts” output from one-hot encoded to integer type; need integer format to obtain the accuracies and to report the matrix



```
> testout = np.argmax(tsLbl, axis=1)
```

Do the same conversion for the test labels

```
> labelname = ['お O', 'き Ki', 'す Su', 'つ Tsu', 'な Na',
               'は Ha', 'ま Ma', 'や Ya', 'れ Re', 'を Wo']
```

```
> testScores = metrics.accuracy_score(testout, predout)
> confusion = metrics.confusion_matrix(testout, predout)
```

Calculates the test scores in accuracy

Calculates the confusion matrix

# Kuzushiji MNIST

## 6. Test model, part 3

- Test the model, calculate the accuracy and confusion matrix

```
> print("Best accuracy (on testing dataset): %.2f%%" % (testScores*100))
> print(metrics.classification_report(testout, predout, target_names=labelname, digits=4))
> print(confusion)
```

```
Best accuracy (on testing dataset): 96.56%
      precision    recall  f1-score   support

   お  o      0.9615      0.9740      0.9677       1000
   き  ki      0.9772      0.9430      0.9598       1000
   す  su      0.9562      0.9390      0.9475       1000
   つ  tsu      0.9732      0.9820      0.9776       1000
   な  na      0.9588      0.9530      0.9559       1000
   は  ha      0.9707      0.9600      0.9653       1000
   ま  ma      0.9245      0.9920      0.9571       1000
   や  ya      0.9877      0.9620      0.9747       1000
   れ  re      0.9665      0.9800      0.9732       1000
   を  wo      0.9838      0.9710      0.9774       1000

avg / total      0.9660      0.9656      0.9656      10000
```

## Classification Report

```
[[974   2   1   1  18   1   0   1   1   1]
 [  5 943   6   0   5   2  24   3   7   5]
 [  8   3 939   9   4   7  19   4   7   0]
 [  0   0   8 982   0   4   5   0   1   0]
 [ 12   2   1   9 953   4   8   2   6   3]
 [  1   3  13   4   1 960  13   0   3   2]
 [  0   2   3   0   1   2 992   0   0   0]
 [  7   5   5   0   5   2   5 962   5   4]
 [  2   1   4   3   5   3   1   0 980   1]
 [  4   4   2   1   2   4   6   2   4 971]]
```

## Confusion Matrix

# Kuzushiji MNIST

## 6. Test model, part 4

Use pandas to read the training log in csv (i.e. modelname + '.csv') that we have saved just now


### • Plot the result

```
> import pandas as pd

records = pd.read_csv(modelname + '.csv')
> plt.figure()
> plt.subplot(211)
> plt.plot(records['val_loss'])
> plt.yticks([0.00,0.10,0.20,0.30])
> plt.title('Loss value',fontsize=12)

> ax = plt.gca()
> ax.set_xticklabels([])

> plt.subplot(212)
> plt.plot(records['val_acc'])
> plt.yticks([0.93,0.95,0.97,0.99])
> plt.title('Accuracy',fontsize=12)
> plt.show()
```



# Project:

Try the different variations in your model and observe their performances

**Rmb to enable your GPU in your COLAB!**

1. The original model; given previously; this will serve as the base for the rest
2. Add 1x CNN and 1x MaxPooling layer; decide on the channels and sizes for these layers
3. Step 2 plus add 2 more Dropout layers of 20%; decide where to put them
4. Steps 2 & 3 plus add an additional dense layer of activation RELU; decide on the number of neurons yourself for this layer
5. A model of your own configuration that gives an optimal accuracy (must be more than the original model's accuracy); **this version to be submitted**



# Project (Optional; only if you have time):

Try the different variations in your model and observe their performances

## Rmb to enable your GPU in your COLAB!

6. From the original model, add padding for all CNN layers
7. From the original model, add 3x additional CNN layers; decide on the channels and sizes for all CNN layers
8. From the original model, add 3x additional dense layers of activation RELU; decide on the number of neurons yourself for all dense layers
9. Steps 7 and 8 together WITHOUT any additional Dropout layers
10. Steps 7 and 8 together WITH additional Dropout layers of 20%; decide how many Dropout layers you need and where to put them

**For those who have finished your quiz but have not finished your workshop, please do so and upload.**

**For those who have finished your quiz and have uploaded your workshop, please give me a while to mark your quiz and check your submission. After which, you may leave early.**