

NUS-ISS

*Pattern Recognition using
Machine Learning System*



Module 8 - Case studies on using recurrent neural networks for machine learning systems

by Dr. Tan Jen Hong

© 2021 National University of Singapore.
All Rights Reserved.

ECG Signals

Normal and Coronary Artery Disease

- Lead II ECG from open source PhysioNet database
- Sampling frequency: 257 Hz
- 5 seconds long, consisting of 1285 data points

Normal



CAD




Source: <https://www.sciencedirect.com/science/article/pii/S0010482517304201>

ECG Signals


Normal and Coronary Artery Disease

- The pathology of coronary artery disease


Normal artery




Endothelial dysfunction




Fatty streak formation



Stable (fibrous) plaque formation



Unstable plaque formation

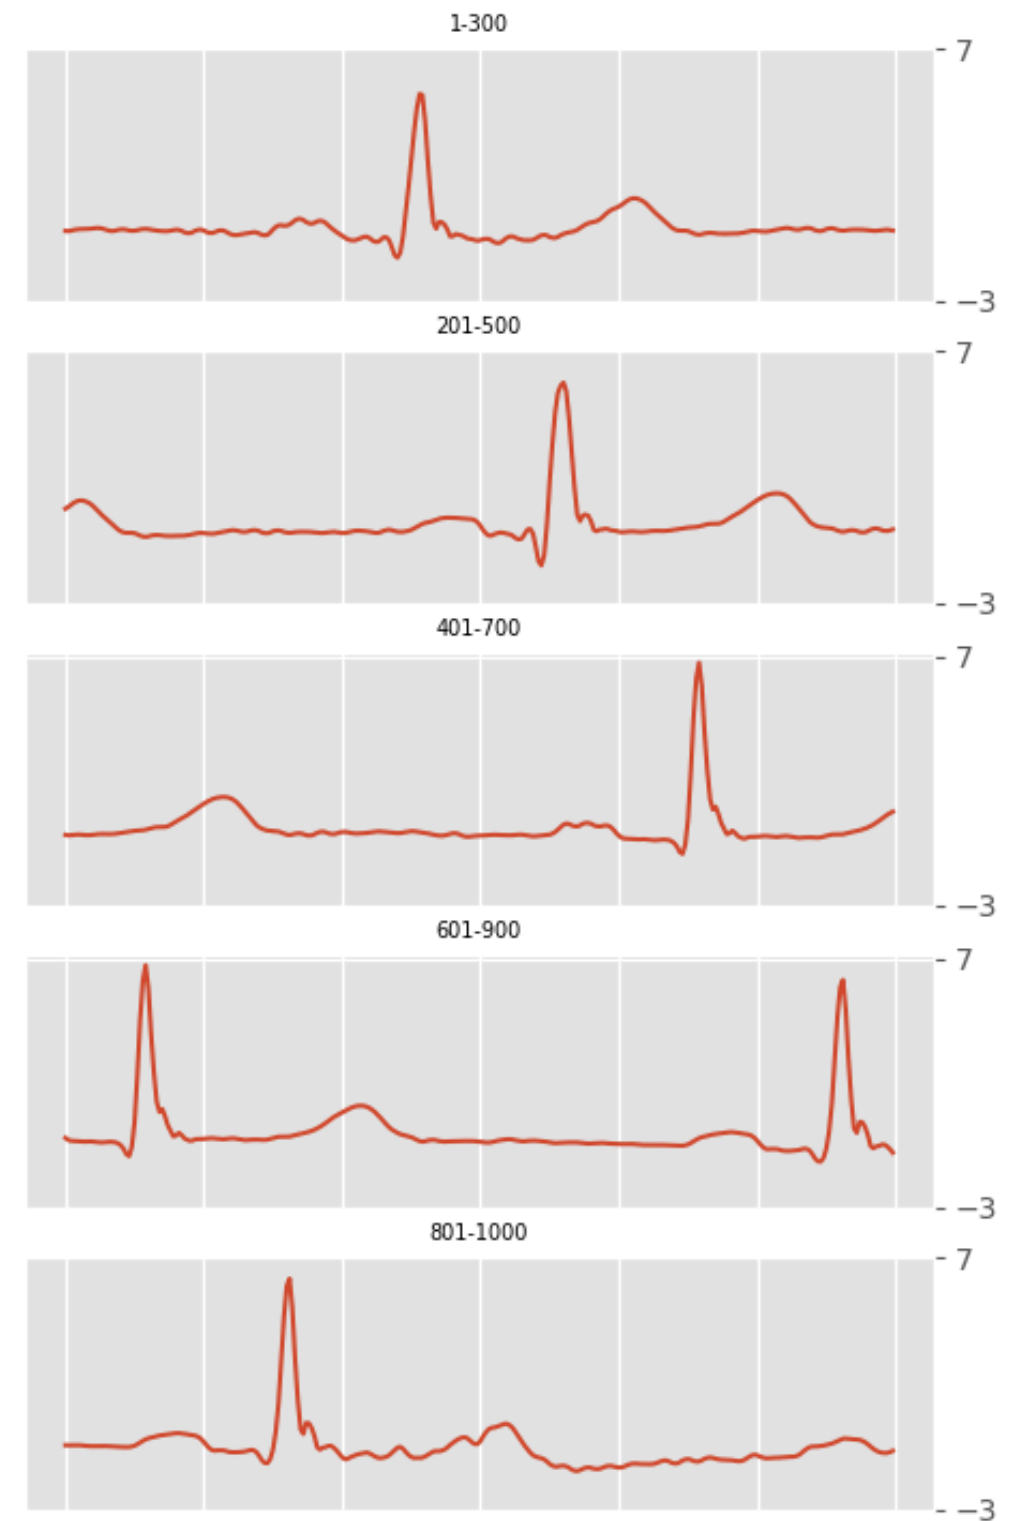


Source: <https://www.sciencedirect.com/science/article/abs/pii/S0950705117302769>

Normal ECG

in segments


length: 300
distance: 200



CAD ECG

in segments

length: 300
distance: 200



ECG

1. Import libraries

- Import all the things that we are going to use in this problem
- h5py is required to load ECG data

```
> import h5py
> import numpy as np
> import sklearn.metrics as metrics
> import matplotlib.pyplot as plt

> from tensorflow.keras.callbacks import ModelCheckpoint, CSVLogger
> from tensorflow.keras.models import Model
> from tensorflow.keras.layers import Dense
> from tensorflow.keras.layers import Input
> from tensorflow.keras.layers import LSTM
```

HDF5

Something about h5py ...

- h5py is the pythonic interface to the HDF5 binary data format
- HDF5: Hierarchical Data Format, designed to store and organize large amounts of data
- File extension: .h5, .hdf5
- Can think HDF as a file system within a file, let you organize data hierarchically
- Keras stores model in this format



Andrew Collette

ECG

2. Matplotlib setup

- Use 'ggplot' style to plot our training and testing result
- The setup uses 'ggplot' style for plot
- Also, for y axis, the labels and ticks put on right rather than left

```
> plt.style.use('ggplot')
> plt.rcParams['ytick.right'] = True
> plt.rcParams['ytick.labelright'] = True
> plt.rcParams['ytick.left'] = False
> plt.rcParams['ytick.labelleft'] = False
> plt.rcParams['font.family'] = 'Arial'
```


ECG

3. Data preparation, part 1

- Load the hdf5 file and extract the data within

```
> f = h5py.File('cad5sec.mat')  
> X = f["data"]  
> Y = f["classLabel"]
```

- if we type 'f' in ipython console, it will return

```
<HDF5 file "cad5sec.mat" (mode r+)>
```

- if we type 'X' and 'Y' respectively in ipython console, it will give

```
<HDF5 dataset "data": shape (1285, 38120), type "<f8">  
<HDF5 dataset "classLabel": shape (1, 38120), type "<f8">
```

ECG

3. Data preparation, part 1

Name ▲	Type	Size
data	float64	(38120, 1285)
label	float64	(38120, 1)

Name ▲	Type	Size
cad	float64	(6120, 1285)
nor	float64	(32000, 1285)

- Need to convert 'X' and 'Y' into numpy array and do a transpose on the data

```
> data      = np.array(X)
> data      = np.transpose(data)
> label     = np.array(Y)
> label     = np.transpose(label)
```

- The first 32000 rows are 'normal' ECG, the rest are 'CAD' ECG

```
> nor       = data[0:32000]
> cad       = data[32000:38120]
```

ECG

3. Data preparation, part 2

- Create the function to turn each sample into fixed number of segments (controlled by length and distance)

```
> def makeSteps(dat, length, dist):  
    width          = dat.shape[1]  
    numOfSteps     = int(np.floor((width-length)/dist)+1)  
  
    segments       = np.zeros([dat.shape[0],numOfSteps,length],  
                               dtype=dat.dtype) ] pre-allocate the array  
  
    for l in range(numOfSteps):  
        segments[:,l,:] = dat[:,(l*dist):(l*dist+length)]  
  
    return segments
```

ECG

3. Data preparation, part 3

- Manually split the data into training and testing set

```
> print('Create dataset...')
> trNor      = nor[0:28800].copy()           (28800, 1285)
> tsNor      = nor[28800:32000].copy()       (3200, 1285)
> trCad      = cad[0:5000].copy()            (5000, 1285)
> tsCad      = cad[5000:6120].copy()         (1120, 1285)
```

- Set the length and the distance, and convert each sample into segments

```
> length     = 36
> dist       = 24

> print('Finalizing all the data.....')
> trNorS     = makeSteps(trNor, length, dist) (28800, 53, 36)
> tsNorS     = makeSteps(tsNor, length, dist) (3200, 53, 36)
> trCadS     = makeSteps(trCad, length, dist) (5000, 53, 36)
> tsCadS     = makeSteps(tsCad, length, dist) (1120, 53, 36)
```

3. Data preparation, part 3

```

> trDat      = np.vstack([trNorS, trCadS])
> tsDat      = np.vstack([tsNorS, tsCadS])

> trLbl      = np.vstack([np.zeros([trNorS.shape[0], 1]),
                           np.ones([trCadS.shape[0], 1])])
> tsLbl      = np.vstack([np.zeros([tsNorS.shape[0], 1]),
                           np.ones([tsCadS.shape[0], 1])])

```

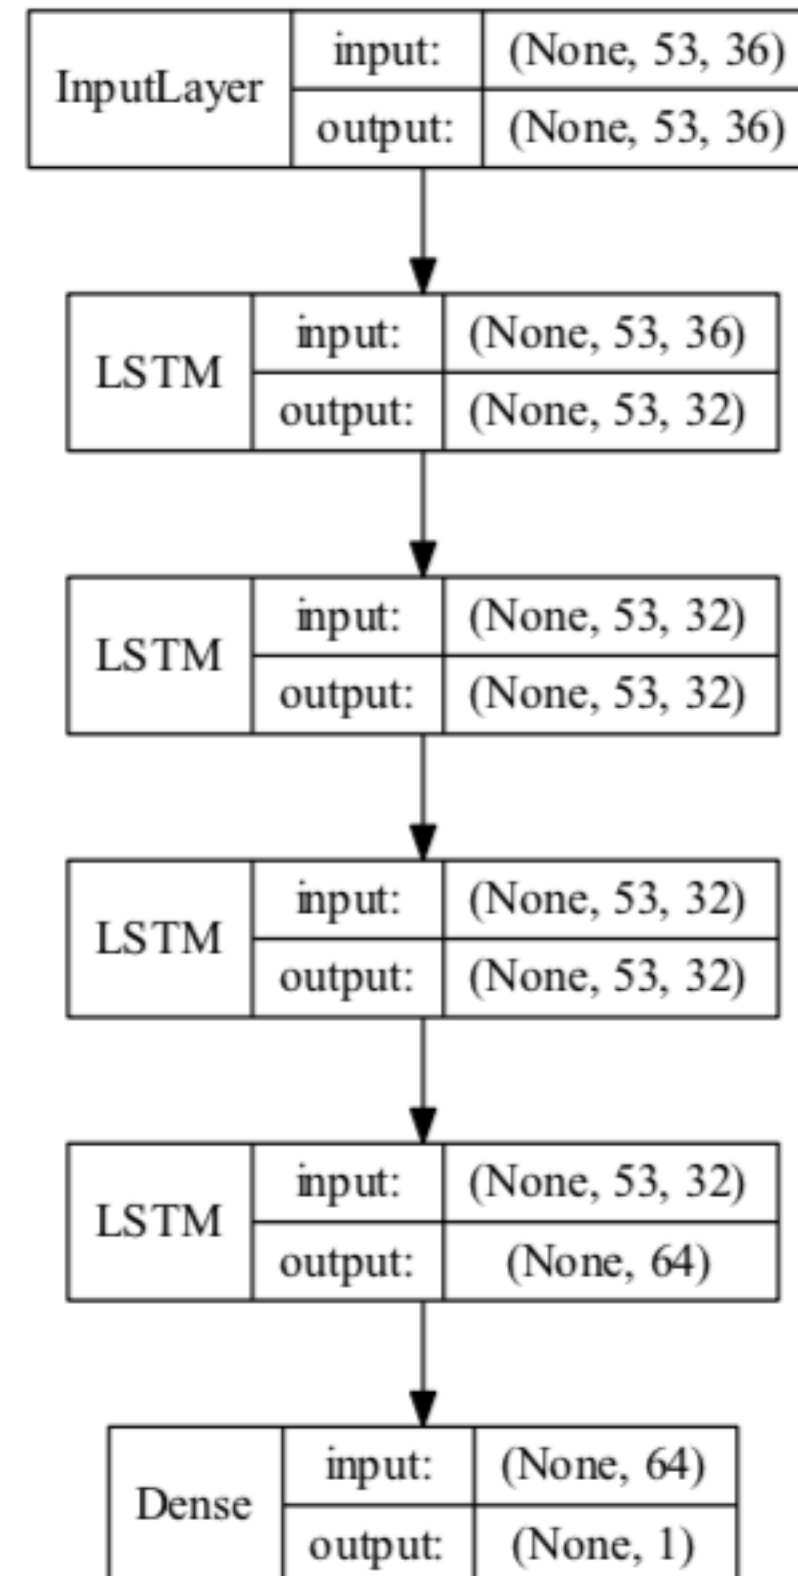
'Normal' set to 0
'Cad' set to 1

trDat	float64	(33800, 53, 36)
trLbl	float64	(33800, 1)

tsDat	float64	(4320, 53, 36)
tsLbl	float64	(4320, 1)

The net

4. Define model



The net

4. Define model

```
> seed = 7
> np.random.seed(seed)
> modelname = 'wks3_2_1'
> def createModel():
    inputs= Input(shape=(trDat.shape[1], length))
    y      = LSTM(32,
                  return_sequences=True,
                  dropout=0.25,
                  recurrent_dropout=0.25)(inputs)
    y      = LSTM(32,
                  return_sequences=True,
                  dropout=0.5,
                  recurrent_dropout=0.25)(y)
    y      = LSTM(32,
                  return_sequences=True,
                  dropout=0.25)(y)
    y      = LSTM(64, dropout=0.25)(y)
    y      = Dense(1, activation='sigmoid')(y)

    model = Model(inputs=inputs, outputs=y)
    model.compile(loss='binary_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])

    return model
```

- Note 1: The dropout layer doesn't work for LSTM / RNN. If want to do dropout, set value to the 'dropout' argument
- Note 2: We can apply dropout on the recurrent part in LSTM (the U_0)
- Note 3: We use 'binary_crossentropy', not 'categorical_crossentropy', because we have only 1 output, and the value is either 1 or 0

The net

4. Define model

- 'model' for training; 'modelGo' for final evaluation

```
> model = createModel()  
> modelGo = createModel()  
  
> model.summary()
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 53, 36)	0
lstm (LSTM)	(None, 53, 32)	8832
lstm_1 (LSTM)	(None, 53, 32)	8320
lstm_2 (LSTM)	(None, 53, 32)	8320
lstm_3 (LSTM)	(None, 64)	24832
dense (Dense)	(None, 1)	65
Total params: 50,369		
Trainable params: 50,369		
Non-trainable params: 0		

The net

4. Define model

- Create checkpoints to save model during training and save training data into csv
- 'monitor' can be 'val_acc' or 'val_loss'
- When set to 'val_acc', 'mode' must be 'max'; when set to 'val_loss', 'mode' must be 'min'

```
> filepath      = modelname + ".hdf5"
> checkpoint    = ModelCheckpoint(filepath,
                                monitor='val_acc',
                                verbose=0,
                                save_best_only=True,
                                mode='max')
```

```
> csv_logger    = CSVLogger(modelname + '.csv')
> callbacks_list = [checkpoint, csv_logger]
```

The net

5. Train model

- The line for training

```
> model.fit(trDat,
            trLbl,
            validation_data=(tsDat, tsLbl),
            epochs=40,
            batch_size=128,
            shuffle=True,
            callbacks=callbacks_list)
```

Train on 33800 samples, validate on 4320 samples

Epoch 1/40

33800/33800 [=====] - 100s 3ms/sample - loss: 0.2614 - acc: 0.9047 - val_loss: 1.0719 - val_acc: 0.7426

Epoch 2/40

33800/33800 [=====] - 96s 3ms/sample - loss: 0.1154 - acc: 0.9616 - val_loss: 0.7840 - val_acc: 0.7681

Epoch 3/40

33800/33800 [=====] - 96s 3ms/sample - loss: 0.0912 - acc: 0.9705 - val_loss: 0.7976 - val_acc: 0.7898

Epoch 4/40

33800/33800 [=====] - 97s 3ms/sample - loss: 0.0797 - acc: 0.9750 - val_loss: 0.6873 - val_acc: 0.8343

Epoch 5/40

33800/33800 [=====] - 97s 3ms/sample - loss: 0.0716 - acc: 0.9772 - val_loss: 0.8630 - val_acc: 0.8188

.....

The net

6. Result


- Classification result

Best accuracy (on testing dataset): 88.38%

	precision	recall	f1-score	support
Normal	0.8766	0.9812	0.9260	3200
CAD	0.9187	0.6054	0.7298	1120
avg / total	0.8875	0.8838	0.8751	4320

- Confusion matrix

```
[[3140  60]
 [ 442 678]]
```



Feature extraction

Better learning

- In previous examples, we simply feed the original data (although with a bit of re-arrangement) into the net
- Feature extraction is missing!
- Without feature extraction, learning is slower and not optimal
- Should perform rounds of feature extraction, reduce the amount of recurrent calculation so that training is faster

Conv1D

size 3

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z	AA
1																											
2		Input							Kernel							Intermediate Output							Output				
3																											
4		1	0	1	0				0	1	0	1				7							7	7			
5		1	1	3	2				0	0	2	0				6							6	13			
6		1	1	0	1				0	1	0	0				8							8	10			
7		2	3	2	1																						
8		0	2	0	1																						
9																											
10		1	0	1	0				2	1	0	0				7											
11		1	1	3	2				0	0	0	1				13											
12		1	1	0	1				0	3	0	0				10											
13		2	3	2	1																						
14		0	2	0	1																						
15																											
16																											

Kernel size (in one direction): 3

Number of filters: 2

Conv 1D

Determine the output size

- M_r, M_c : Number of rows and columns in the output, respectively
- W_r : Number of rows in the input
- F_r : Filter size (along the rows)
- P_r : Amount of padding (along the rows)
- S_r : Stride (along the rows)
- N_F : The number of filters

$$M_r = \left\lfloor \frac{W_r + P_r - F_r}{S_r} \right\rfloor + 1$$

$$M_c = N_F$$

Conv 1D

Calculate the parameters

- For each filter, the number of parameters is

$$W_c \times F_r + 1$$

- Thus, for N_F amount of filters, the number of parameters is

$$N_F \times (W_c \times F_r + 1)$$

M_r, M_c : Number of rows and columns in the output, respectively

W_c : Number of columns in the input

F_r : Filter size (along the rows)

N_F : The number of filters

MaxPooling1D

size 3

	A	B	C	D	E	F	G	H	I	J	K	L	M	N
1														
2		Input							Output					
3														
4		1	0	1	0				1	1	3	2		
5		0	1	3	2				2	3	2	1		
6		1	1	2	1									
7		2	3	0	0									
8		0	2	0	1									
9														
10														

Kernel size (in one direction): 2
 Stride (in one direction): 2

Stacked Conv1D LSTM

ECG classification in Keras

```
> def createModel():
    inputs = Input(shape=(trDat.shape[1],length))
    y = Conv1D(32, 5, activation='relu')(inputs)
    y = Dropout(0.25)(y)
    y = Conv1D(32, 5, activation='relu')(y)
    y = MaxPooling1D(2)(y)
    y = Conv1D(48, 5, activation='relu')(y)
    y = Dropout(0.5)(y)
    y = Conv1D(48, 5, activation='relu')(y)
    y = MaxPooling1D(2)(y)
    y = Conv1D(64, 5, activation='relu')(y)
    y = Dropout(0.5)(y)
    y = Conv1D(64, 5, activation='relu')(y)
    y = MaxPooling1D(2)(y)
    y = LSTM(8,
              return_sequences=True,
              dropout=0.5,
              recurrent_dropout=0.5)(y)
    y = LSTM(4,
              return_sequences=True,
              dropout=0.5,
              recurrent_dropout=0.5)(y)
    y = LSTM(2)(y)
    y = Dense(1, activation='sigmoid')(y)

    model = Model(inputs=inputs,outputs=y)
    model.compile(loss='binary_crossentropy',
                  optimizer='adam',
                  metrics=['accuracy'])

    return model
```

length = 24
dist = 6

Stacked Conv1D LSTM

ECG classification in Keras

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 211, 24)	0
conv1d (Conv1D)	(None, 207, 32)	3872
dropout (Dropout)	(None, 207, 32)	0
conv1d_1 (Conv1D)	(None, 203, 32)	5152
max_pooling1d (MaxPooling1D)	(None, 101, 32)	0
conv1d_2 (Conv1D)	(None, 97, 48)	7728
dropout_1 (Dropout)	(None, 97, 48)	0
conv1d_3 (Conv1D)	(None, 93, 48)	11568
max_pooling1d_1 (MaxPooling1D)	(None, 46, 48)	0
conv1d_4 (Conv1D)	(None, 42, 64)	15424
dropout_2 (Dropout)	(None, 42, 64)	0
conv1d_5 (Conv1D)	(None, 38, 64)	20544
max_pooling1d_2 (MaxPooling1D)	(None, 19, 64)	0
lstm (LSTM)	(None, 19, 8)	2336
lstm_1 (LSTM)	(None, 19, 4)	208
lstm_2 (LSTM)	(None, 2)	56
dense (Dense)	(None, 1)	3
Total params: 66,891		
Trainable params: 66,891		

Stacked Conv1D LSTM

The result


- Classification result

Best accuracy (on testing dataset): 96.44%

	precision	recall	f1-score	support
Normal	0.9997	0.9522	0.9754	3200
CAD	0.8797	0.9991	0.9356	1120
avg / total	0.9686	0.9644	0.9651	4320

- Confusion matrix

[[3047 153]
[1 1119]]



— training
— testing