



SENSE MAKING OF SENSOR DATA

Dr TIAN Jing

tianjing@nus.edu.sg



Module objective

Module: Making sense of sensor data

Knowledge and understanding

- Understand the fundamentals of sense making pipeline of both single type of sensor data and multiple types of sensor data

Key skills

- Design, build, implement intelligent sensing system for real-world applications


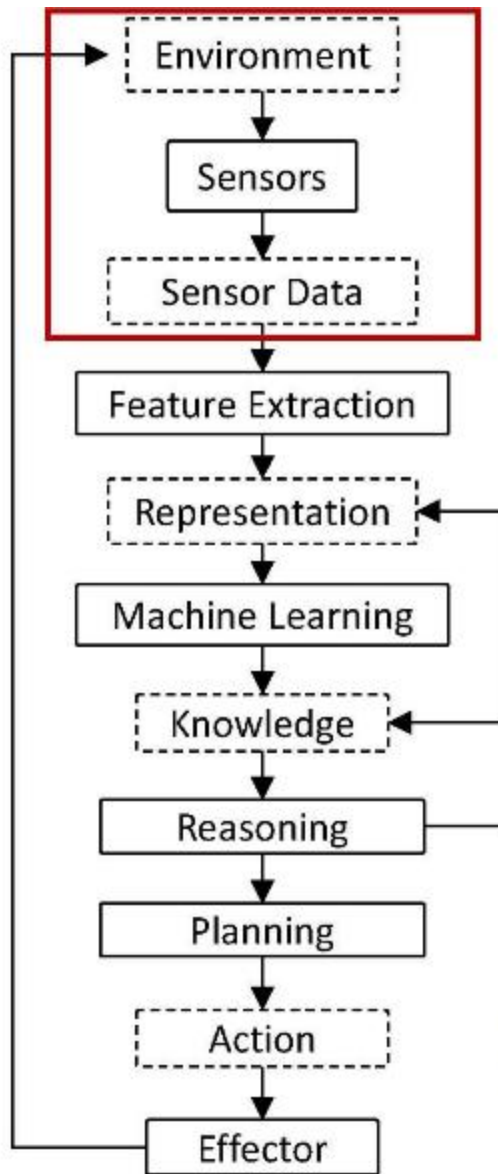


Major reference

- [Introduction] O. Tickoo and R. Iyer, ***Making Sense of Sensors: End-to-End Algorithms and Infrastructure Design from Wearable-Devices to Data Centers***, Apress, 2016.
- [Intermediate] L. A. Klein, ***Sensor and Data Fusion: A Tool for Information Assessment and Decision Making***, SPIE Press, 2012.
- [Advanced] T. Giannakopoulos, ***Multimodal Information Processing & Analysis***,
<https://github.com/tyiannak/multimodalAnalysis>
- [Advanced] ***Advanced Multimodal Machine Learning***, CMU,
<https://piazza.com/cmu/spring2017/11777/home>



Overview of intelligent sensing system (1)



Lidar



Camera
(Visible, Infrared)



Radar



GPS



Stereo Camera



Microphone



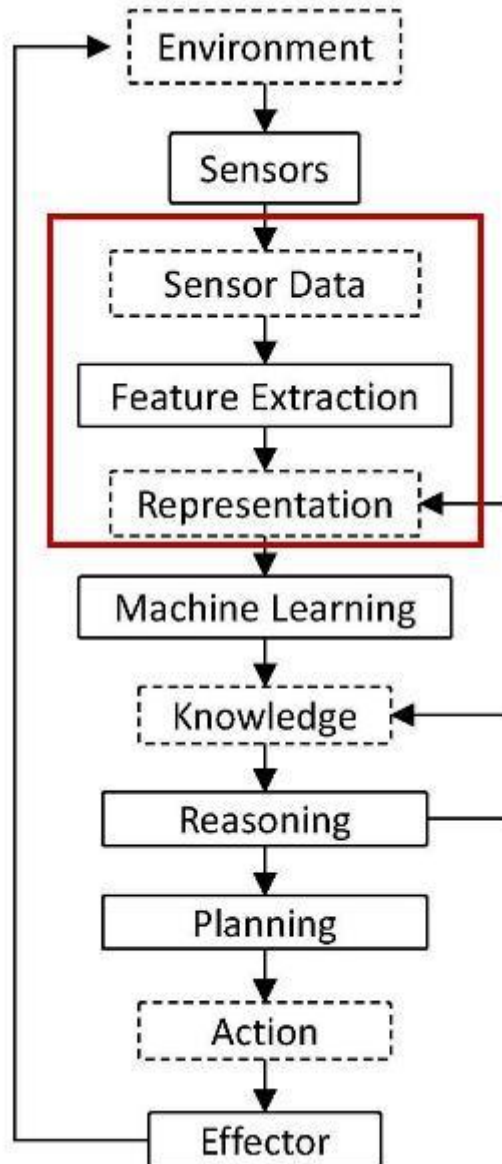
Networking
(Wired, Wireless)



IMU



Overview of intelligent sensing system (2)



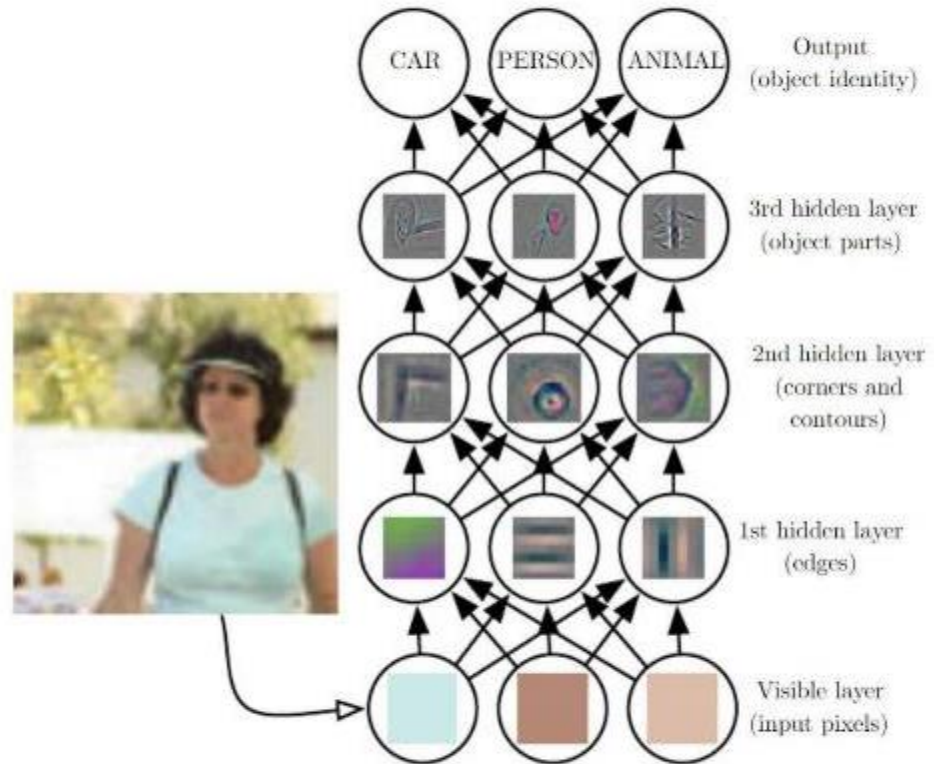
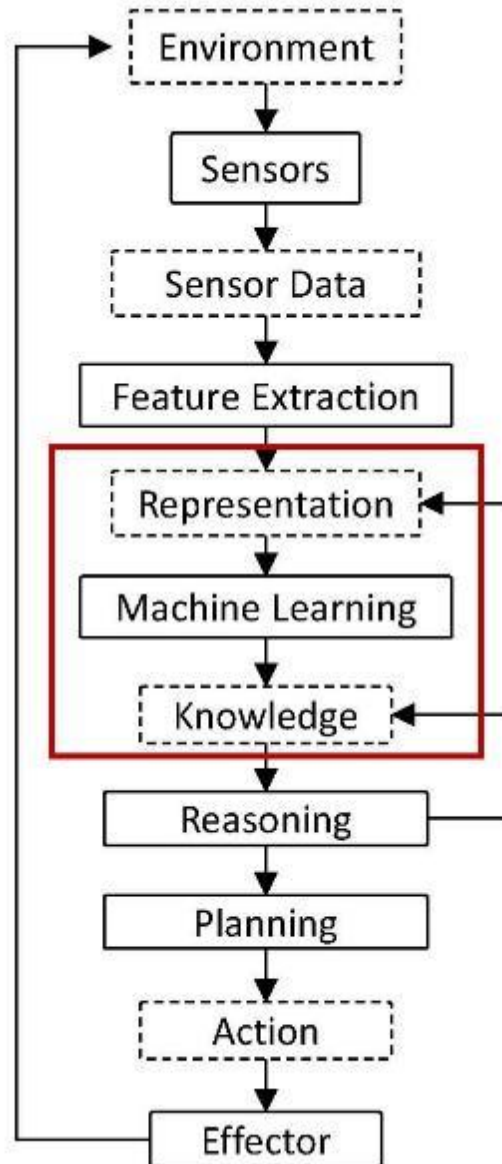
1. Signal representation can be **manually designed** (input-agnostic)

2. Signal representation can be **adaptively to the input signal**

3. Signal representation can be **learned from signal dataset**



Overview of intelligent sensing system (3)



Source: MIT 6.S099: Artificial General Intelligence, <https://agi.mit.edu/>



Overview of intelligent sensing system (4)

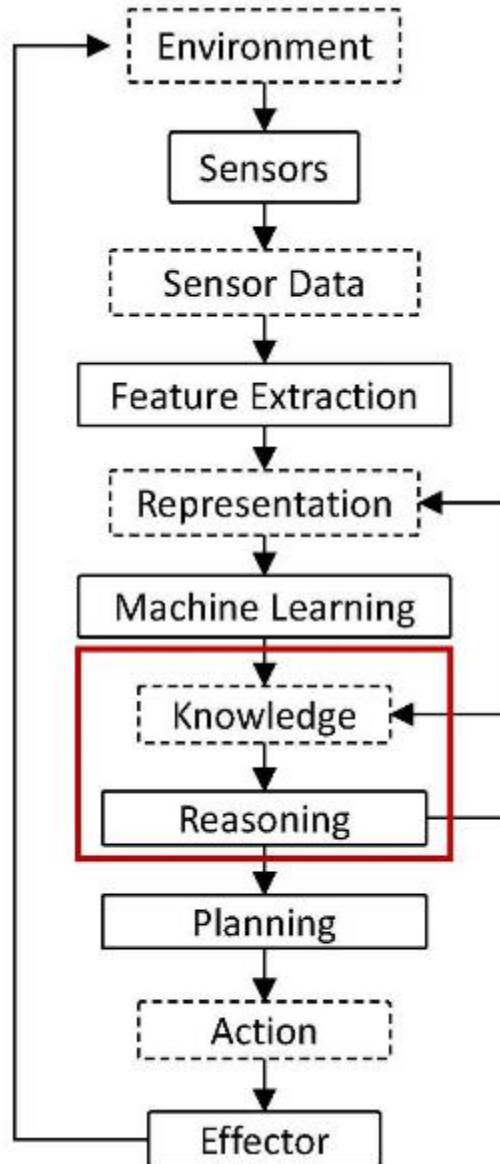


Image Recognition:
If it looks like a duck



Audio Recognition:
Quacks like a duck

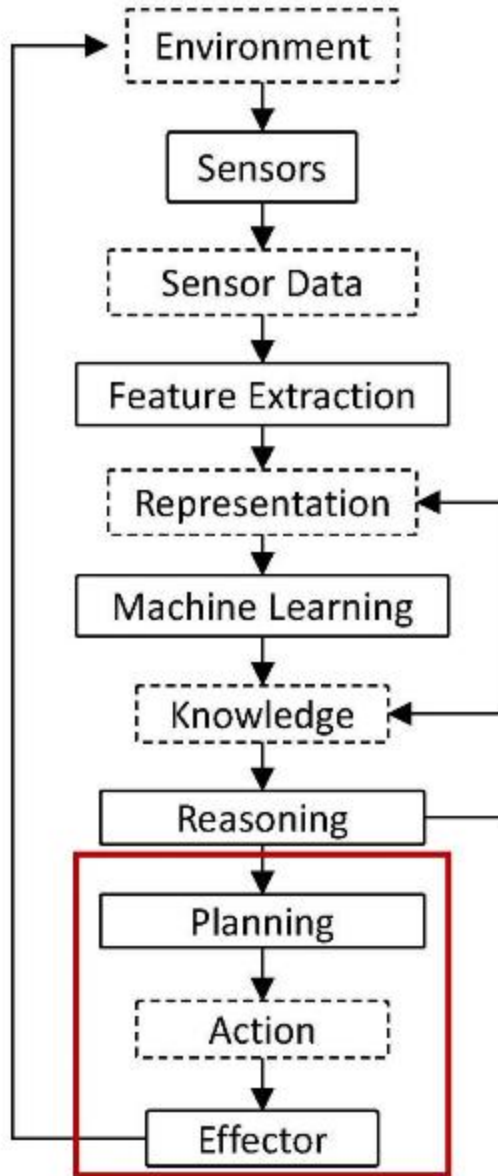


Activity Recognition:
Swims like a duck





Overview of intelligent sensing system (5)

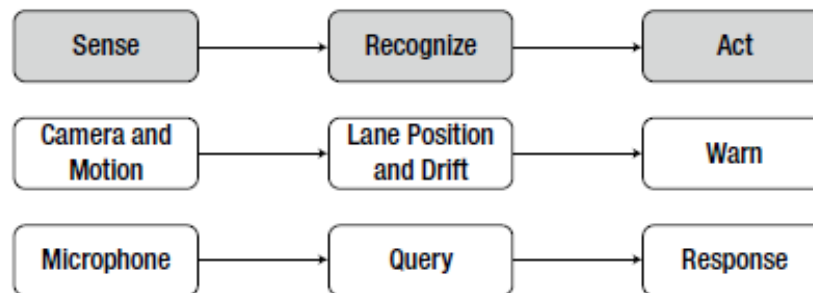




An example

Example: Automatic Driver Assistance System (ADAS)

- It uses a set of sensors like cameras, motion sensors, etc. to provide a driver with assistance for safety and comfort, such as tracking the position of a car in the traffic lanes and provide audible warnings in case the car drifts from its lane (Lane Departure Warning).
- It works on the fundamental processing principles of *Internet of Things* (IoT) comprising of sensors collecting data (cameras, motion sensors), algorithms recognizing the data (relative position of the car in the lane, car drift), and applications acting on the recognized data (audible warning on/off).



Source: O. Tickoo and R. Iyer, Making Sense of Sensors: End-to-End Algorithms and Infrastructure Design from Wearable-Devices to Data Centers, Apress, 2016.



Making sense of sensor data pipeline (1)



Raw Sensor Data

- Raw (or unprocessed) sensor data is captured by sensors at the front end of the pipeline
- **Vision:** Camera sensors take still or video shots of a scene
- **Audio/Speech:** The microphone-like sensors typically capture the audio and analog-to-digital conversion is applied
- **Motion:** Inertial sensors typically measure the acceleration and motion
- **BMI sensors:** The Brain Machine Interface sensors, like EEG, report the brain activity measurements as activity graphs
- **Other IoT sensors:** temperature, humidity, etc.




Making sense of sensor data pipeline (2)



Category	Representative features
Geometrical	Edges, lines, line widths, line relationships (e.g., parallel, perpendicular), circles, shapes, size of enclosed area
Structural	Surface area; relative orientation; orientation in vertical and horizontal ground plane
Statistical	Number of surfaces, area and perimeter, mean, variance, kurtosis, skewness, entropy
Spectral	Color coefficients, spectral peaks and lines
Time domain	Pulse characteristics (rise and fall times, amplitude), pulse width, pulse repetition interval, moments
Frequency domain	Fourier coefficients, Wavelet coefficients



Making sense of sensor data pipeline (3)




Recognized Entities

- Complex task involving spatial and temporal analysis of the extracted features to map the aggregate to pre-known entities.
- For vision, the recognition task may involve classifying the extracted features to recognize shapes like objects and faces.
- For audio, statistical analysis and classification of features allows for aggregated features to be recognized as words.



Making sense of sensor data pipeline (4)



Semantic Relationship

- The connection between various recognized entities, for example, the semantic relationship between a key and a lock refers to the operation of opening or closing the lock with the key.
- For audio, the sentence "*I like coffee*" connects "I" (the subject) with "coffee" (object) using the predicate "like."
- For vision, semantic context of the video can be analyzed based on identified entities (multiple objects, locations, people, and activities) and scoring these entities with respect to their co-occurrence as well as relation to the type of classification scenario of importance.



Making sense of sensor data pipeline (5)



Contextual Relationship

- Understand the recognized data in context. The context can be obtained in multiple ways. These include using other sensors as well as using the history of recognized data for context recognition.
- For vision: A face recognition surveillance system can identify a potential invasion or normal operation depending on when the recognition takes place (day vs. night) or depending on the person that is recognized.



Making sense of sensor data pipeline (6)



Knowledge

- Knowledge representation: Knowledge needs to be stored and represented in a manner such that the semantic information and relationships between various concepts is retained and is modifiable.
- Knowledge operations: These include tools, APIs, and programming languages to insert and extract data from the knowledge database. This involves understanding the incoming semantic data and providing a mechanism to find and manipulate the correct relationships. The extraction process involves responding to different queries targeting relationships between different entities and concepts.




Single type of sensor data: Inertial

- An accelerometer essentially measures the force (proper acceleration) along x, y and z-axes.

Use cases


- Understanding of position/orientation helps mobile phones re-orient the screen in portrait or landscape mode and reverse direction as required.
- Gesture recognition based on buffering continuous data and looking at the change in force and orientation.



Source: <http://www.utdallas.edu/~kehtar/UTD-MHAD.html>



Example: Berkeley Multimodal Human Action Database (MHAD)



Source: http://tele-immersion.citris-uc.org/berkeley_mhad



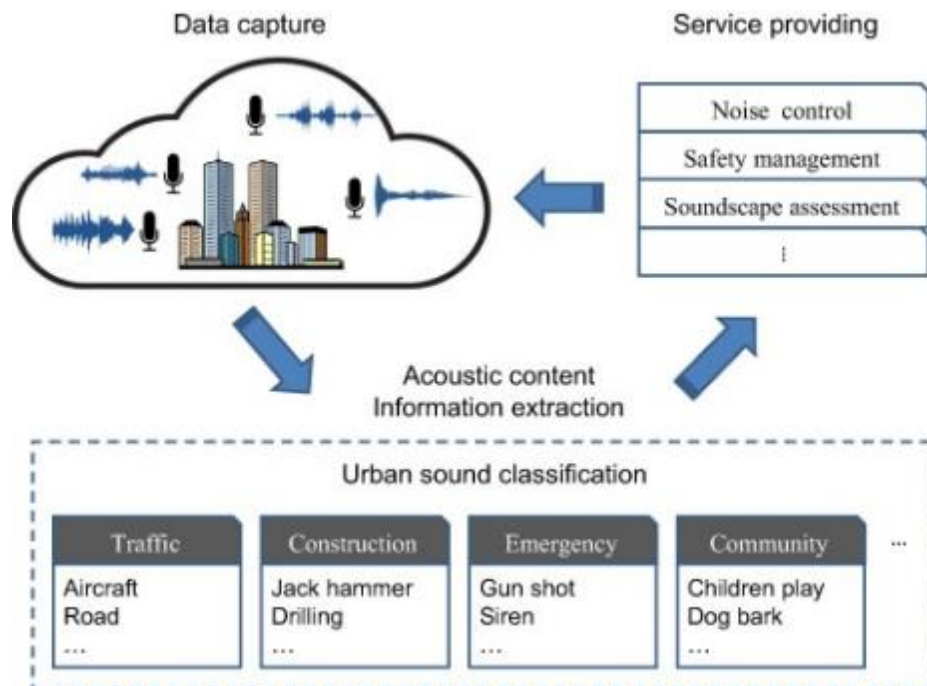
Single type of sensor data: Audio

- **Audio classification:** A common IoT use case for microphones is to classify the environment that the sound was captured in.
- **Voice activity detection:** Determine whether there is voice in the captured audio.
- **Speaker recognition:** Speaker recognition, voice recognition, attempts to determine who is speaking.
- **Keyword recognition:** Recognize whether a particular word was uttered. Keyword recognition can also be generalized to keyphrase recognition and both are typically used as triggers for additional activity such as starting a session of commands or bringing up an application.
- **Command and control:** Refers to using a small set of phrases in speech recognition. For illustration, this could include a set of commands to control a toy car such as “move forward,” “move backward,” “go faster,” “go slower,” “turn right/left,” etc.



Example: Urban sound classification

- Dataset: Urban sound dataset, <https://urbansounddataset.weebly.com/urbansound.html>
- It contains 1302 labeled sound recordings. Each recording is labeled with the start and end times of sound events from 10 classes: air_conditioner, car_horn, children_playing, dog_bark, drilling, engine_idling, gun_shot, jackhammer, siren, and street_music.



Source: <https://www.sciencedirect.com/science/article/abs/pii/S0003682X16302274>



Single type of sensor data: Vision

- **Scene recognition:** Identify multiple objects, faces, and people in an image and using that information to determine the likely activity or context.
- **Object recognition:** Identify objects in an image and potentially matching them to a pre-existing database of objects that have been captured before.
- **Gesture recognition:** Recognize static poses or moving gestures either specific to the hand/arm or the human body.
- **Anomaly detection:** Identify if any anomaly occurred which should trigger additional analysis.
- **Video summarization:** Summarize the salient aspects of a long video stream. This includes scene changes, key scenarios and objects/characters that are the focus of the video.

Example: Video understanding

- Dataset: YouTube-8M, a video understanding challenge, <https://ai.googleblog.com/2017/02/an-updated-youtube-8m-video.html>



Source: <https://ai.googleblog.com/2017/02/an-updated-youtube-8m-video.html>



Sensor data fusion (1)

- **Representation:** The information of the fusion process has an abstract level higher than each input data set.
- **Certainty:** If V is the sensor data before fusion and $p(V)$ is the *a priori* probability of the data before fusion, then the gain in certainty is the growth in $p(V)$ after fusion.
- **Accuracy:** The standard deviation on the data after the fusion process is smaller than the standard deviation provided directly by the sources.
- **Completeness:** Bringing new information to the current knowledge on an environment allows a more complete the view on this environment.



Sensor data fusion (2)

- **Fusion across sensors (same type):** A number of temperature sensors measuring the temperature of an object.
- **Fusion across domains (same type, different definition/range):** A number of sensors measure the same attribute over a number of different ranges or domains. This arises, for example, in the definition of a temperature scale.
- **Fusion across attributes (different types):** A number of sensors measure different quantities associated with the same experimental situation, such as measuring air temperature, pressure and humidity to determine air refractive index.
- **Fusion across time:** Current measurements are fused with historical information.

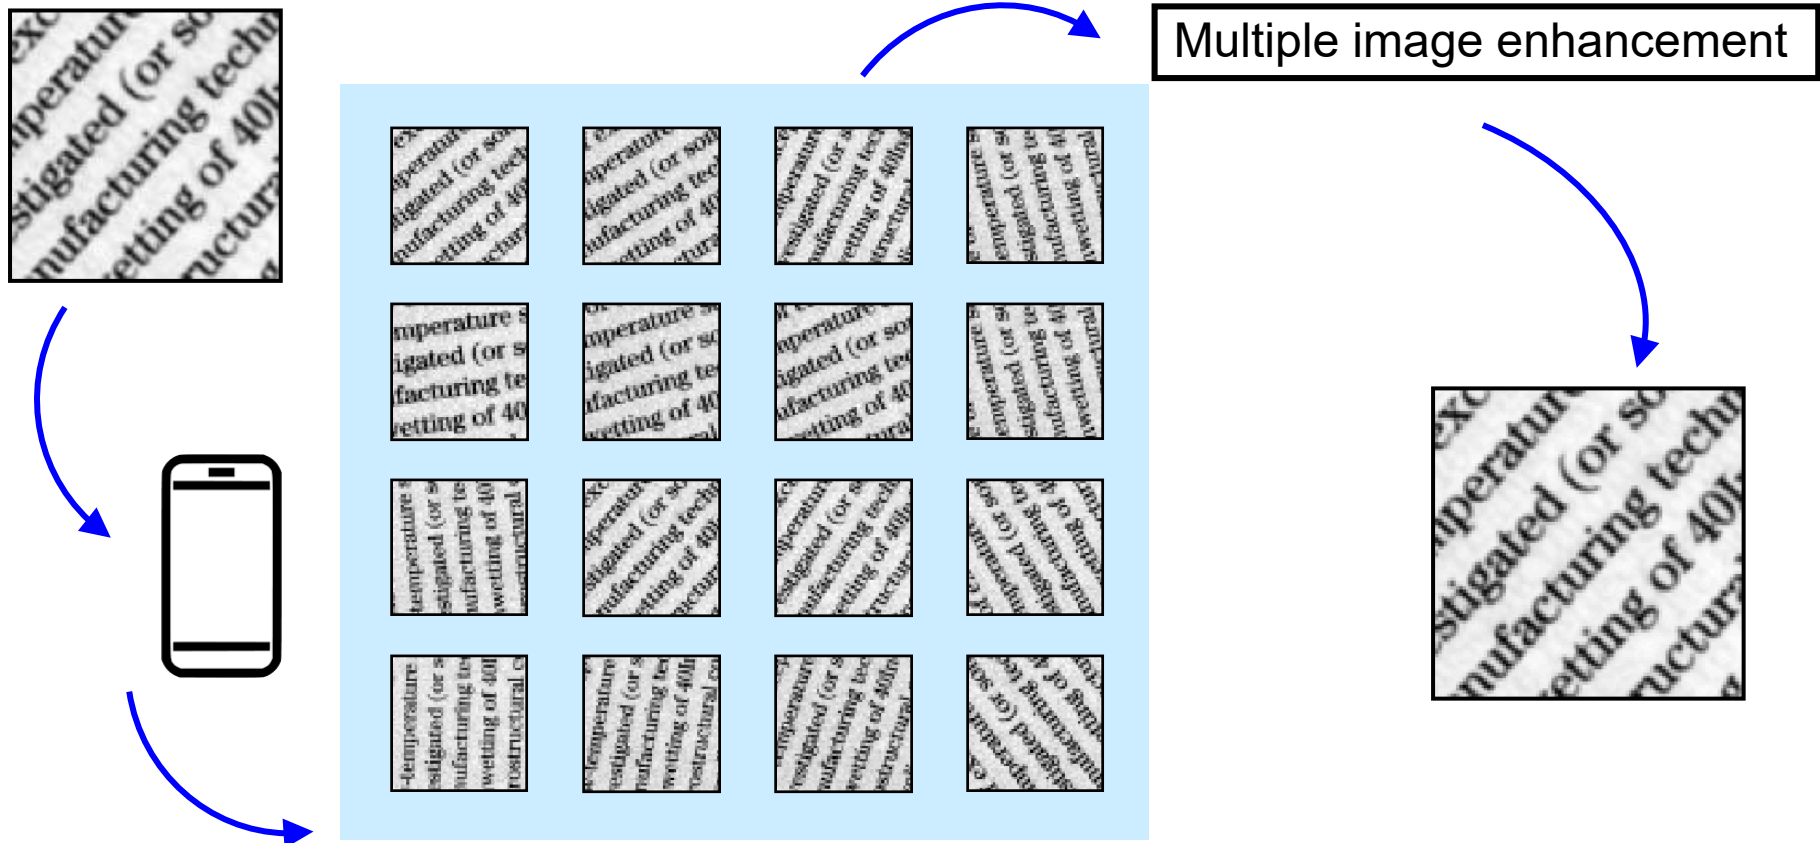


Sensor data registration

- **Spatial alignment.** Transformation of the local spatial positions to a common coordinate system. The process involves geo-referencing the location and field-of-view of each sensor.
- **Temporal alignment.** Transformation of the local times to a common time axis. In many applications, is performed using a dynamic time warping algorithm.




Sensor data registration (1)

- At different times (multi-temporal fusion)





Sensor data registration (2)

- With different sensors (multi-modal fusion)



Sensor data registration (3)

- From different viewpoints (multi-view fusion)







Sensor data registration (4): Audio

Euclidean registration (top figure): Align the i -th point on one time series with the same i -th point.

Dynamic time warping registration (bottom figure): Allows similar shapes to match even if they are different in the time axis.



Reference: <https://databricks.com/blog/2019/04/30/understanding-dynamic-time-warping.html>



Sensor data registration (5): Video

Objective: Find point to point correspondence between two frames in the video sequence.

The pixel intensities of an object do not change between consecutive frames. Neighbouring pixels have similar motion.





Photo: <http://users.umiacs.umd.edu/~zhuolin/Keckgesturedataset.html>



Data augmentation

- Data augmentation can be viewed as an injection of prior knowledge about the invariant properties of the data against certain transformations. Augmented data can cover unexplored input space, prevent overfitting, and improve the generalization ability of a machine learning model.
- **Scaling** changes the magnitude of the data in a window by multiplying by a random scalar.
- **Magnitude filtering** changes the magnitude of each sample by convolving the data window with a smooth curve varying around one.
- **Jittering** is simulating additive sensor noise. These data augmentation methods may increase robustness against multiplicative and additive noise and improve performance.
- **Time-warping** perturbs the temporal location. By smoothly distorting the time intervals between samples, the temporal locations of the samples can be changed using time-warping.

Reference:



- T. T. Um, "Data augmentation of wearable sensor data for parkinson's disease monitoring using convolutional neural networks," ACM Int. Conf. on Multimodal Interaction, New York, NY, USA, 2017, pp. 216–220.
- Time Series Data Augmentation for Deep Learning: A Survey, <https://arxiv.org/pdf/2002.12478.pdf>



Data augmentation



Scaling

Magnitude filtering




Jittering

Time-warping




- **Uncoupled sensor data fusion / Late fusion:** The data fusion occurs at the last possible stage in the respective pipelines. The advantage of this method is that the sensor fusion is simple and can be accomplished using existing technology pipelines for recognition. Domain experts like this approach because cross-domain technical knowledge requirements are minimal.






Multimodal sensor data recognition (2)

- **Loosely coupled sensor data fusion / Restricted fusion:**
Restricts the recognition search space of mode modality based on the results from the other. In this approach, one recognition pipe helps to set the context for the other. Once that is accomplished, the second recognition pipe need only perform recognition within the established context boundaries, leading to potential savings in compute and higher performance.



- **Tightly coupled sensor data fusion / Early fusion:** It allows for combination of multiple sensors that appear as a single complex sensor for most of the applications. Data representation at the lower levels of stack is very sensor-specific, so experts are needed to make sensible decisions for merging data from multiple sources. Each sensor model has a different data representation format at the lowest level, making it very hard to come up with a homogenous representation that preserves the information content.



- Making sense of sensor data pipeline
- Single type of sensor data analytics
- Multiple types of sensor data analytics

Thank you!

Dr TIAN Jing
Email: tianjing@nus.edu.sg