



EB5004 PLP – NEW MEDIA AND SENTIMENT MINING

ENTITY & ASPECT MINING

Dr. Fan Zhenzhen
Institute of Systems Science
National University of Singapore

© NUS. The contents contained in this document may not be reproduced in any form or by any means, without the written permission of ISS, NUS, other than for the purpose for which it has been supplied.



Module Objectives

At the end of this module, you will be able to

- To extract key entities and aspects from opinions for closer analysis of sentiments and their targets
- To learn the main methods of entity/aspect extraction

- The opinion target - entity and its aspects
- Tasks in ABSA
- Entity extraction
- Aspect extraction
- Recent advances



The Opinion Target and its Aspects

- An *opinion* consists of a *sentiment* (positive or negative) and a *target* (of opinion).
- Detect subjective opinion and determine its polarity is often the first step in sentiment mining.
- A common follow-up question is : what exactly do people like or not like?
- Getting the right answer requires **finer-grained** analysis, identifying opinion *target*, which can be an *entity*, or its *aspects*.
- Crucial for applications like product review analysis.



Definition of Opinion

- Recall that opinion is formally defined as a quadruple (g, s, h, t) , where
 - g is the sentiment target,
 - s is the sentiment of the opinion about the target g ,
 - h is the opinion holder (the person or organization who holds the opinion),
 - and t is the time when the opinion is expressed..
- In the case of entity and aspect mining, there is an additional dimension – aspect/ feature **a** .
- Extricate quintuple (e, a, s, h, t) with entity e and aspect **a** together representing the opinion target.



Opinion Target Extraction

- Examples:
 - *Although the **service** is **not that great**, I still **love** this **restaurant**.*
 - *The **iPhone**'s **call quality** is **good**, but its **battery life** is **short**.*
- **Entities** - names of products, services, individuals, events, and organizations
- **Aspects** – the attributes and components of entities
- Identifying them from text data is called *opinion target extraction*.
- Information Extraction (IE) tasks.




Typical tasks for aspect

- Identify subjective/opinionated sentences
- Find and extract entities and aspects that have been commented on by an opinion holder
- Determine whether the opinions are positive, negative or neutral
- Group entity/aspect synonyms
- Produce an aspect-based opinion summary of multiple reviews



ABSA subtasks



Subtask	Input	Output	Task Type
Aspect Term Extraction(<i>AE</i>)	S	a_1, a_2	Extraction
Opinion Term Extraction(<i>OE</i>)	S	o_1, o_2	Extraction
Aspect-level Sentiment Classification(<i>ALSC</i>)	$S + a_1$ $S + a_2$	s_1 s_2	Classification
Aspect-oriented Opinion Extraction(<i>AOE</i>)	$S + a_1$ $S + a_2$	o_1 o_2	Extraction
Aspect Term Extraction and Sentiment Classification(<i>AESC</i>)	S	(a_1, s_1) , (a_2, s_2)	Extraction & Classification
Pair Extraction(<i>Pair</i>)	S	(a_1, o_1) , (a_2, o_2)	Extraction
Triplet Extraction(<i>Triplet</i>)	S	(a_1, o_1, s_1) , (a_2, o_2, s_2)	Extraction & Classification

- The methods and features used for recognizing **entities** and **aspects** are usually different due to their individual specific characteristics.
- Main idea:
 - An opinion always has a target.
 - The target is an aspect or an entity.
 - It's usually a noun or noun phrase (but not always).
 - => often leverage on syntactic structures to discover opinion and target relationships



ENTITY EXTRACTION



Entity Extraction

- Similar to the classic problem of **Named Entity Recognition** (NER) in NLP.
- Main approaches
 - Rule-based
 - Supervised statistical machine learning (e.g., HMMs, CRFs, etc.)
 - Semi-supervised approaches that only require some unambiguous seed entity names to find all entities of the same type (e.g., PU learning, Bayesian Sets, etc)
- Mature tools are available to perform NER (GATE, NLTK, Stanford NER, etc.), but not for entity extraction.
- For sentiment analysis, an additional step is usually needed to group named entities into synonyms, **entity resolution** (ER), as the entity can be mentioned in various ways.



Differences from traditional NER

- In traditional NER, the objective is to recognize all named entities of certain types in a corpus, e.g., names of people, names of organizations, etc.
- In sentiment mining, the interest is often to find the mentions of a set of desired entities, e.g. the company's own products, its competitors' products, etc.
- Solved in two steps:
 - Identify all entity mentions or entity expressions of the interested set of entities.
 - For each entity mention/expression, determine to which entity it belongs, aka *entity linking* or *entity disambiguation*, which is a special case of ER



Entity Resolution

- Two name ambiguity problems:
 - Polysemy – *Apple* for *Apple Inc.* (the maker of iPhone and iPad), or *Apple Daily* (a Hong Kong newspaper)?
 - Synonymy – *National University of Singapore* vs. *NUS*, *Volkswagen* vs. *Vwagen*, *Singapore* vs. *the little red dot*, *sound/voice/sound quality* etc.
- When the set of target entities is available, the task is to cluster entity mentions, and map them to target entities.



Entities in different types of text data

- Entity-focused corpora – online reviews of products and services.
 - Entity information can be obtained from meta-data
 - Mentions of other entities might need to be identified for comparison opinion mining
- Domain-focused corpora – forum discussions.
 - Normally focusing on discussions of a specific type of products or topics.
 - Need to perform entity extraction and linking
- Open domain corpora – like Twitter.
 - Can contain documents of **any** entity or topic, with little or no meta-data (like hashtags)
 - Obviously entity extraction and linking is needed. Most challenging.



Search helps

- In a very large corpus, keyword search is often applied first to retrieve relevant posts, using name variations of the desired entities.
- Thus, it's better to have a clear mining objective, from which you'll derive a comprehensive list of name variations to search the large corpus, like Twitter.
- Due to polysemy, a filtering step may be required to get posts that truly contain the desired entities.



How: Semi-Supervised Extraction

- Supervised entity extraction using HMM or CRF is still applicable here, but it requires labeling of training data, which may not always be possible.
- Therefore the semi-supervised approaches are often applied; e.g. PU learning (learning from positive and unlabeled examples), Bayesian sets
- Given a set of seed entity names, identify all entities of the same type as the seeds from a given corpus
- General idea: if they appear in similar context, they are likely to be the same kind!



- For example, to identify entities of phones
 - Given a set of positive examples (of seed entities), and a corpus
 - Find candidate entities from the corpus
 - Using sequences of specific POS tags: NNP (proper noun), NNPS (plural proper noun), and CD (cardinal number)
 - E.g. “**Samsung/NNP Galaxy/NNP S5/CD**” as a candidate
 - For each seed entity, create a TF vector for each mention of it representing a positive example, using the surrounding words context of the seed mention.
 - Do the same for each candidate entity
 - Use a PU learning algorithm like S-EM (Liu et al., 2002) to learn a classification model to label the candidates.

Li, X., L. Zhang, B. Liu, and S. Ng. Distributional similarity vs. PU learning for entity set expansion. In *Proceedings of Annual Meeting of the Association for Computational Linguistics (ACL-2010)*, 2010b.



How: Supervised Entity Linking

- To identify if a mention refers to some known entity, or none.
 - But instead of PER(person), ORG(organization), GPE(geopolitical entity), UKN(unknown), we are typically looking for **products, services, and brands**.
- What's required: the set of desired entities, each with its disambiguating text and type (KB); and entity expressions to be linked
- Example solution, supervised:
 - Candidate generation: for each entity expression, generate possible entities from KB using heuristic rules -> multiple (expression, entity) pairs
 - Candidate ranking: learning to rank the pairs (features such as similarity of entity name strings, similarity of context, entity type, etc.)
 - Or classification: learn a classifier to predict if the (expression, entity) pair is positive/negative
 - Case of NIL (no applicable entity from KB): when no pair is predicted to be positive.
- Main challenges: feature engineering, and labeled examples



Another complexity

- In sentiment analysis applications involving consumer products, products typically have brands and models, which can form a hierarchical relationship.
- Brand – product – model
 - e.g. Apple – iPhone – iPhone X
- Usually by separating brands and models



Opinion Holder and Time Extraction

- NER task!
- For social media data
 - the opinion holder is usually the author of the review, blog, post, etc.
 - It's trivial when the ID and date/time of the post are usually known.
- For other data, like news articles, they may need to be extracted out of text.
 - E.g. consider person and organization entities, and score them using Maximum Entropy model.



ASPECT EXTRACTION



Main approaches

1. By finding frequent nouns and noun phrases.
2. By exploiting syntactic relations
 - Syntactic dependencies depicting opinion and target relations
 - Lexico-syntactic patterns encoding entity and part/attribute relations
3. Using supervised learning



1. Frequency-based approach

- Assumption: a reasonable number of reviews about the same product or at least about the same type of products.
 - Find nouns and noun phrases using a POS tagger
 - Count their occurrence frequencies
 - Keep only the frequent ones above a experimentally determined threshold
- **It works.**
 - Aspects are usually expressed as nouns and noun phrases.
 - Vocabulary converges when people comment on the same (type of) product.
 - Irrelevant contents tend to be infrequent.



Frequency-based aspect extraction

- Frequency is the key here.
- Simple and effective. The candidate aspects are almost always the most important aspects of the product.
- Applicable to entity extraction too.
- Caution: Won't work if the corpus has a mixture of very different products and/or if each product has only one or two reviews.



Better precision with co-occurrence

- To improve its precision, use heuristics to filter for noun phrases that are more likely to be aspects of entities.
 - E.g. for camera reviews, phrases indicating *part-of* relations, like “*of camera*”, “*camera has*”, “*camera comes with*”, can be used to find camera components by web search
 - The discovered phrases that often **co-occur** with such *part-of* (*meronymy*) relation indicators are likely to be correct aspects.
 - A simplified version of Point-wise Mutual Information (PMI) can be used to compute the co-occurrence strength

$$PMI(a, d) = \frac{hits(a \wedge d)}{hits(a)hits(d)}$$



Parts or Attributes

- To distinguish components/parts from attributes
- Can use Morphological cues (e.g., “-*iness*,” “-*ity*” suffixes)
- WordNet is a great resource
 - Synonyms (“*fast*” and “*quick*”), antonyms (“*fast*” and “*slow*”)
 - Hyponyms and hypernyms (*is-a* hierarchy), meronyms and holonyms (*part-whole* relation)
 - For example, we can query WordNet with the word “*camera*” and get its senses

Noun

- **S: (n)** camera, [photographic camera](#) (equipment for taking photographs (usually consisting of a lightproof box with a lens at one end and light-sensitive film at the other))
 - [direct hyponym](#) / [full hyponym](#)
 - [part meronym](#)
 - [direct hypernym](#) / [inherited hypernym](#) / [sister term](#)
- **S: (n)** [television camera](#), [tv camera](#), **camera** (television equipment consisting of a lens system that focuses an image on a photosensitive mosaic that is scanned by an electron beam)

<https://wordnet.princeton.edu/>

- Hyponyms (x is-a camera)

- direct hyponym / full hyponym

- S: (n) box camera, box Kodak (a simple camera shaped like a rectangular box)
 - S: (n) candid camera (a miniature camera with a fast lens)
 - S: (n) digital camera (a camera that encodes an image digitally and store it for later reproduction)
 - S: (n) webcam (a digital camera designed to take digital photographs and transmit them over the internet)
 - S: (n) flash camera (a camera with a photoflash attachment)
 - S: (n) motion-picture camera, movie camera, cine-camera (a camera that takes a sequence of photographs that can give the illusion of motion when viewed in rapid succession)
 - S: (n) sound camera (a movie camera that records sounds in synchrony with the visual images)
 - S: (n) point-and-shoot camera (a lightweight photographic camera with an autofocus)
 - S: (n) Polaroid camera, Polaroid Land camera (a camera that develops and produces a positive print within seconds)
 - S: (n) portrait camera (a camera with a portrait lens)
 - S: (n) reflex camera (camera that allows the photographer to view and focus the exact scene being photographed)

- Hypernyms (camera is-a x)

- direct hypernym / inherited hypernym / sister term

- S: (n) photographic equipment (equipment used by a photographer)
 - S: (n) equipment (an instrumentality needed for an undertaking or to perform a service)
 - S: (n) instrumentality, instrumentation (an artifact (or system of artifacts) that is instrumental in accomplishing some end)
 - S: (n) artifact, artefact (a man-made object taken as a whole)

- S: (n) whole, unit (an assemblage of parts that is regarded as a single entity) *"how big is that part compared to the whole?"*; *"the team is a unit"*

- S: (n) object, physical object (a tangible and visible entity; an entity that can cast a shadow) *"it was full of rackets, balls and other objects"*

- S: (n) physical entity (an entity that has physical existence)

- S: (n) entity (that which is perceived or known or inferred to have its own distinct existence (living or nonliving))

- direct hypernym / inherited hypernym / sister term

- S: (n) photographic equipment (equipment used by a photographer)
 - S: (n) camera, photographic camera (equipment for taking photographs (usually consisting of a lightproof box with a lens at one end and light-sensitive film at the other))
 - S: (n) clapperboard (photographic equipment used to synchronize sound and motion picture; boards held in front of a movie camera are banged together)
 - S: (n) developer (photographic equipment consisting of a chemical solution for developing film)
 - S: (n) enlarger (photographic equipment consisting of an optical projector used to enlarge a photograph)
 - S: (n) flash, photoflash, flash lamp, flashgun, flashbulb, flash bulb (a lamp for providing momentary light to take a photograph)
 - S: (n) light meter, exposure meter, photometer (photographic equipment that measures the intensity of light)
 - S: (n) photographic paper, photographic material (light-sensitive paper on which photograph can be printed)
 - S: (n) sensitometer (a measuring instrument for measuring the light sensitivity of film over a range of exposures)

- Meronyms (x is part-of camera)

- part meronym

- S: (n) aperture (a device that controls amount of light admitted)
- S: (n) camera lens, optical lens (a lens that focuses the image in a camera)
- S: (n) delayed action (a mechanism that automatically delays the release of a camera shutter for a fixed period of time so that the photographer can appear in the picture)
- S: (n) diaphragm, stop (a mechanical device in a camera that controls size of aperture of the lens) *"the new cameras adjust the diaphragm automatically"*
- S: (n) finder, viewfinder, view finder (optical device that helps a user to find the target of interest)
- S: (n) hood, lens hood (a tubular attachment used to keep stray light out of the lens of a camera)
- S: (n) magazine, cartridge (a light-tight supply chamber holding the film and supplying it for exposure as required)
- S: (n) shutter (a mechanical device on a camera that opens and closes to control the time of a photographic exposure)
- S: (n) sprocket (roller that has teeth on the rims to pull film or paper through)



Further refinements

- Filter by dropping aspects that do not have sufficient mentions alongside known sentiment words (Blair-Goldensohn et al., 2008)
- Collapse aspects at the word stem level
- Pattern-based filter to remove non-aspect expressions (Moghaddam and Ester, 2010)
- Compare the frequencies of the candidates in a review corpus with those in a generic corpus to identify true aspects (Scaffidi et al., 2007)
- First find the core aspect words, then use the information distance to find other related words, e.g. “\$” and “*dollars*” for “*price*” (Long et al., 2007)
- Etc., etc.



2. Exploiting syntactic relations

- Syntactic relations between sentiment expressions and their sentiment or opinion targets (Needs a parser)
 - E.g. “*This camera takes great photos.*”
 - “*Picture quality and battery life are great.*”
- Linguistic constructions
 - E.g. X of Y: “the *voice quality* of the *iPhone*”
 - Genitives: “the *camera’s price*”



Opinion-target relations

- Can use a dependency parser
 - Identify reliable dependency relation templates from training(labelled) data
 - Then use them to find valid aspect-sentiment pairs in test data
 - Can be used for simultaneous extraction of both sentiment words and opinion targets – Double Propagation (DP) method (Qiu et al. 2009, 2011)
 - E.g. “The *software* is *amazing*.” -> pattern “NN – nsubj –JJ”

Tagging

The/DT software/NN is/VBZ amazing/JJ ./.

Parse

```
(ROOT
  (S
    (NP (DT The) (NN software))
    (VP (VBZ is)
      (ADJP (JJ amazing)))
    (. .)))
```



Universal dependencies, enhanced

```
det(software-2, The-1)
nsubj(amazing-4, software-2)
cop(amazing-4, is-3)
root(ROOT-0, amazing-4)
```

Stanford Parser
<http://nlp.stanford.edu:8080/parser/index.jsp>



Dependencies

- Dependencies are binary relations: a grammatical relation between
 - A *governor* (also known as *head*), and
 - A *dependent*
- For example, *amod* (adjectival modifier)
 - An adjectival modifier of an NP is any adjectival phrase that serves to modify the meaning of the NP.
 - E.g. *I like this amazing software.*



Universal Dependencies

	Nominals	Clauses	Modifier words	Function Words
Core arguments	<u>nsubj</u> <u>obj</u> <u>iobj</u>	<u>csubj</u> <u>ccomp</u> <u>xcomp</u>		
Non-core dependents	<u>obl</u> <u>vocative</u> <u>expl</u> <u>dislocated</u>	<u>advcl</u>	<u>advmod</u> * <u>discourse</u>	<u>aux</u> <u>cop</u> <u>mark</u>
Nominal dependents	<u>nmod</u> <u>appos</u> <u>nummod</u>	<u>acl</u>	<u>amod</u>	<u>det</u> <u>clf</u> <u>case</u>
Coordination	MWE	Loose	Special	Other
<u>conj</u> <u>cc</u>	<u>fixed</u> <u>flat</u> <u>compound</u>	<u>list</u> <u>parataxis</u>	<u>orphan</u> <u>goeswith</u> <u>reparandum</u>	<u>punct</u> <u>root</u> <u>dep</u>



Stanford Typed Dependencies

- *mod* – modifier

amod - adjectival modifier

appos - appositional modifier

advcl - adverbial clause modifier

det - determiner

predet - predeterminer

preconj - preconjunct

vmod - reduced, non-finite verbal modifier

mwe - multi-word expression modifier

mark - marker (word introducing an *advcl* or *ccomp*)

advmod - adverbial modifier

neg - negation modifier

rcmod - relative clause modifier

quantmod - quantifier modifier

nn - noun compound modifier

npadvmod - noun phrase adverbial modifier

tmod - temporal modifier

num - numeric modifier

number - element of compound number

prep - prepositional modifier

poss - possession modifier

possessive - possessive modifier ('s)

prt - phrasal verb particle



Stanford Typed Dependencies

- *arg* - arguments

agent - agent

comp - complement

acompl - adjectival complement

ccomp - clausal complement with internal subject

xcomp - clausal complement with external subject

obj - object

dobj - direct object

iobj - indirect object

pobj - object of preposition

subj - subject

nsubj - nominal subject

nsubjpass - passive nominal subject

csubj - clausal subject

csubjpass - passive clausal subject



Dependency Relations

- Common dependency relations between sentiment words and aspects include *amod*, *prep*, *nsubj*, *csubj*, *xsubj*, *dobj* and *iobj*
- Common relations for sentiment words and aspects themselves include: the conjunction relation *conj*



More examples

- “*The phone has a nice screen.*”

Tagging

The/DT phone/NN has/VBZ a/DT nice/JJ screen/NN .

Universal dependencies, enhanced

```
det(phone-2, The-1)
nsubj(has-3, phone-2)
root(ROOT-0, has-3)
det(screen-6, a-4)
amod(screen-6, nice-5)
dobj(has-3, screen-6)
```

- “*I like the color of the phone.*”

Tagging

I/PRP like/VBP the/DT color/NN of/IN the/DT phone/NN

Universal dependencies, enhanced

```
nsubj(like-2, I-1)
root(ROOT-0, like-2)
det(color-4, the-3)
dobj(like-2, color-4)
case(phone-7, of-5)
det(phone-7, the-6)
nmod:of(color-4, phone-7)
```



Example Rules for Aspect and Opinion Word Extraction

Rule ID	Observed Relation (Line 1) and Constraints (Lines 2–4)	Output	Examples
R1 ₁ (OA-Rel)	$O \rightarrow O-Dep \rightarrow A$ <i>s.t.</i> $O \in \{O\}$, $O-Dep \in \{MR\}$, $POS(A) \in \{NN\}$	$a = A$	The phone has a <u>good</u> “screen.” <i>good</i> \rightarrow <i>mod</i> \rightarrow <i>screen</i>
R1 ₂ (OA-Rel)	$O \rightarrow O-Dep \rightarrow H \leftarrow A-Dep \leftarrow A$ <i>s.t.</i> $O \in \{O\}$, $O/A-Dep \in$ $\{MR\}$, $POS(A) \in \{NN\}$	$a = A$	“iPod” is the <u>best</u> MP3 player. <i>best</i> \rightarrow <i>mod</i> \rightarrow <i>player</i> \leftarrow <i>subj</i> \leftarrow <i>iPod</i>
R2 ₁ (OA-Rel)	$O \rightarrow O-Dep \rightarrow A$ <i>s.t.</i> $A \in \{A\}$, $O-Dep \in \{MR\}$, $POS(O) \in \{JJ\}$	$o = O$	Same as R1 ₁ with <i>screen</i> as the known word and <i>good</i> as the extracted word
R2 ₂ (OA-Rel)	$O \rightarrow O-Dep \rightarrow H \leftarrow A-Dep \leftarrow A$ <i>s.t.</i> $A \in \{A\}$, $O/A-Dep \in$ $\{MR\}$, $POS(O) \in \{JJ\}$	$o = O$	Same as R1 ₂ with <i>iPod</i> as the known word and <i>best</i> as the extract word
R3 ₁ (AA-Rel)	$A_{i(j)} \rightarrow A_{i(j)}-Dep \rightarrow A_{j(i)}$ <i>s.t.</i> $A_{j(i)} \in \{A\}$, $A_{i(j)}-Dep \in$ $\{CONJ\}$, $POS(A_{i(j)}) \in \{NN\}$	$a = A_{i(j)}$	Does the player play DVDs with <u>audio</u> and “video”? <i>video</i> \rightarrow <i>conj</i> \rightarrow <i>audio</i>
R3 ₂ (AA-Rel)	$A_i \rightarrow A_i-Dep \rightarrow H \leftarrow A_j-Dep \leftarrow A_j$ <i>s.t.</i> $A_i \in \{A\}$, $A_i-Dep = A_j-$ <i>Dep</i> OR ($A_i-Dep =$ <i>subj</i> AND $A_j-Dep =$ <i>obj</i>), POS (A_j) $\in \{NN\}$	$a = A_j$	Canon “G3” has a great <u>lens</u> . <i>len</i> \rightarrow <i>obj</i> \rightarrow <i>has</i> \leftarrow <i>subj</i> \leftarrow <i>G3</i>
R4 ₁ (OO-Rel)	$O_{i(j)} \rightarrow O_{i(j)}-Dep \rightarrow O_{j(i)}$ <i>s.t.</i> $O_{j(i)} \in \{O\}$, $O_{i(j)}-Dep \in$ $\{CONJ\}$, $POS(O_{i(j)}) \in \{JJ\}$	$o = O_{i(j)}$	The camera is <u>amazing</u> and “easy” to use. <i>easy</i> \rightarrow <i>conj</i> \rightarrow <i>amazing</i>
R4 ₂ (OO-Rel)	$O_i \rightarrow O_i-Dep \rightarrow H \leftarrow O_j-$ <i>Dep</i> $\leftarrow O_j$ <i>s.t.</i> $O_i \in \{O\}$, $O_i-Dep = O_j-$ <i>Dep</i> OR (O_i / O_j-Dep $\in \{pnmod, mod\}$), POS (O_j) $\in \{JJ\}$	$o = O_j$	If you want to buy a <u>sexy</u> , “cool,” accessory-available MP3 player, you can choose iPod. <i>sexy</i> \rightarrow <i>mod</i> \rightarrow <i>player</i> \leftarrow <i>mod</i> \leftarrow <i>cool</i>

- Genitive constructions are frequently used to express part-of and attribute-of relations.
 - “the battery *of* the iPhone” – “NP-head of NP-mod”
 - “the iPhone’s sound quality” – “NP-mod’s NP-head”
- However, the semantic relations of the two nouns can be quite different in different context
 - **Part-of:** “iPhone’s battery”
 - Possession: “John’s iPhone”
 - **Attribute-of:** “iPhone’s price”
 - Kinship: “John’s brother”
 - Source-from: “John’s birth city”
 - Make-produce: “Apple’s phone”
- In sentiment analysis, it’s easier to **fix NP-mod to be a specific entity to look for its aspect**



The approach with no parser

- Full parsing is expensive (typically fewer than 20 sentences per second), and doesn't work well on informal data.
- **Approximate** the dependency using distance:
 - If a sentence does not have a frequent aspect but has some sentiment words, the **nearest** noun or noun phrase to a sentiment word is extracted as an aspect.
- **Approximate** the dependency relations
 - Using linear patterns of words and POS tags, or chunk patterns (from shallow parsing)
 - And extract using a good pattern matching algorithm.
- Very useful in practice.



Ranking the candidates

- Aspect relevance ($r(a)$): how likely it's a genuine aspect
 - If an aspect is modified by multiple sentiment words
“*delivery*” modified by “*quick*”, “*cumbersome*”, “*timely*”
 - If it's extracted by multiple lexico-syntactic patterns
 - “*The engine of the car is large*” and “*The car has a big engine*”
 - If it's extracted by both a sentiment word modification relation and a lexico-syntactic pattern in the same sentence
 - “*There is a bad hole in the mattress*”
- Aspect frequency ($f(a)$)
 - Rank the frequent aspects higher
- Final ranking score: $S(a) = r(a) \log (f(a))$



Going further

- Use phrase dependency parser to extract noun phrases and verb phrases.
- Adding comparative- and superlative-based relations
 - E.g. “*The iPhone 5 has better voice quality than Moto X.*”
- Adding sentiment composition rules
 - E.g. “*Enbrel has reduced my joint pain.*”
- Resource usage aspect - resource expression, (*usage_verb, quantifier, resource_noun*)
 - E.g. “*This washer uses a lot of water.*”
usage




3. Supervised learning

- Treat aspect extraction as a special IE problem.
- Dominant method: sequential learning or **sequence labeling** like *hidden Markov models* (HMMs) and *conditional random fields* (CRF).
- Many other methods, like sequential rules, tree-structured classification, etc. have been tried as well.
- Requires labelled data.

- Directed sequence model, successfully applied to many sequence labeling problems such as NER and POS tagging

$\mathbf{y} = \langle y_0, y_1, \dots, y_t \rangle$: hidden state sequence

$\mathbf{x} = \langle x_0, x_1, \dots, x_t \rangle$: observation sequence



- Assumptions:
 - state y_i only depends on its immediate predecessor state y_{i-1} (Markov Assumption)
 - the observation x_i only depends on the current state y_i .
- For aspect extraction,
 - Observations: words or phrases in a review
 - Underlying states: aspect or opinion expression tags

- The **joint probability** of a state sequence \mathbf{y} and an observation sequence \mathbf{x}

$$p(\mathbf{y}, \mathbf{x}) = \prod_{i=1}^t p(y_i | y_{i-1}) p(x_i | y_i)$$

- $p(y_0)$ over the initial state,
- a state transition distribution $p(y_i | y_{i-1})$,
- an observation distribution $p(x_i | y_i)$
- Learning the model: given some observation sequences, learn the model parameter that maximizes the observation probability, e.g. MLE from counts.
- Applying the model (decoding): find an optimal state sequence for a new observation sequence, e.g. Viterbi algorithm



E.g. Lexicalized HMM


- Given a sequence of words $W = w_1 w_2 w_3 \dots w_n$ and corresponding parts-of-speech $S = s_1 s_2 s_3 \dots s_n$,
- The task is to find an appropriate sequence of tags $T = t_1 t_2 t_3 \dots t_n$ that maximize the conditional probability $P(T|W, S)$
- For example: “*I love the ease of transferring the pictures to my computer.*”
- Tags:

```
<BG>I</BG><OPINION_POS_EXP>love</OPINION_P  
OS_EXP><BG>the</BG><PROD_FEAT-  
BOE>ease</PROD_FEAT-BOE> <PROD_FEAT-MOE>  
of</PROD_FEAT-MOE><PROD_FEAT-  
MOE>transferring</PROD_FEAT-MOE>  
<PROD_FEAT-MOE>the</PROD_FEAT-MOE>  
<PROD_FEAT-EOE>pictures</PROD_FEAT-EOE>  
<BG>to</BG><BG>my</BG><BG>computer</BG>
```



Understanding HMM

- Hidden states: H(ot) and C(ol)d weather
- Observations: the number of ice creams eaten on that day – $\{1, 2, 3\}$
- Transition probabilities
- Emission probabilities






Sequence probability

- E.g. to compute the joint probability of observation sequence $\langle 3, 1, 3 \rangle$ and hidden state sequence $\langle \text{hot}, \text{hot}, \text{cold} \rangle$

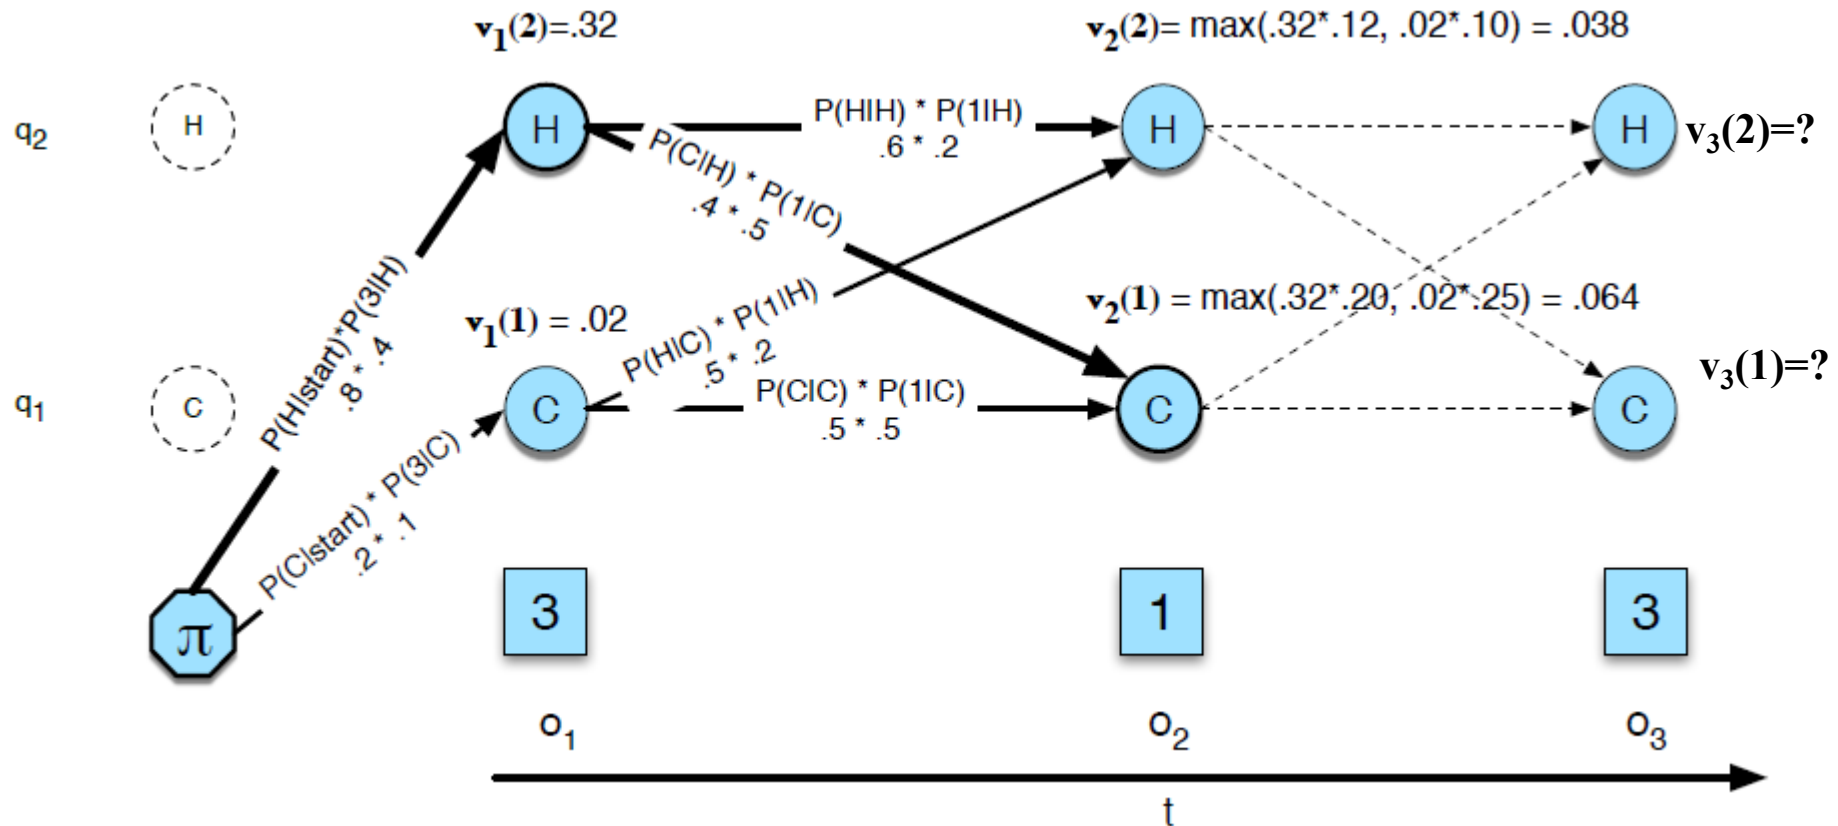
$$P(3 \ 1 \ 3, \text{hot hot cold}) = P(\text{hot}|\text{start}) \times P(\text{hot}|\text{hot}) \times P(\text{cold}|\text{hot}) \\ \times P(3|\text{hot}) \times P(1|\text{hot}) \times P(3|\text{cold})$$





Viterbi Decoding

- Given the observation sequence $\langle 3, 1, 3 \rangle$, find the best sequence of hidden states





Exercise

- The previous slide shows the computation of $v_t(j)$ for two states at two time steps: $v_1(1)$ & $v_1(2)$, $v_2(1)$ & $v_2(2)$

$$v_t(j) = \max_{1 \leq i \leq N-1} v_{t-1}(i) a_{ij} b_j(o_t)$$

- Now try to compute the values for $v_3(1)$ & $v_3(2)$



Conditional Random Fields

- CRF is a discriminative classifier model (HMM is a generative model)
- Extension of logistic regression to sequential data, easier to incorporate a lot of features
- Undirected graphical model, thus its features are not restricted to tags of the preceding nodes -> bidirectional
- It models a conditional probability $p(\mathbf{y}|\mathbf{x})$ over hidden sequence \mathbf{y} given observation sequence \mathbf{x}
- Trained to label an unknown observation sequence \mathbf{x} by selecting the hidden sequence \mathbf{y} that maximizes $p(\mathbf{y}|\mathbf{x})$.



CRF in NLP context

- Suppose (X, Y) is a conditional random field such that X are the '*observables*' and Y is a '*latent*' variable. In NLP, Y can be the EA tags (that you want to learn) while X could be the POS tags.

$$p(\mathbf{y}|\mathbf{x}) = \underbrace{\frac{1}{Z(\mathbf{x})}}_{\text{Normalization}} \prod_{t=1}^T \exp \left\{ \sum_{k=1}^K \underbrace{\theta_k}_{\text{Weight}} \underbrace{f_k(y_t, y_{t-1}, \mathbf{x}_t)}_{\text{Feature}} \right\}$$

- E.g. the EA tag of a word can be determined by the word, its POS tag, its neighbouring words with their respective POS and EA tags
- The expression is a conditional probability which is computed by Bayes rule statistically from the learning corpus and filters out the most likely sequence of X and Y .



Mathematical intuition - CRF

	the	only	redeeming	factor	was	the	food
X	DT	JJ	NN	NN	VBD	DT	NN
Y	O	O	O	O	O	O	B-A

- The CRF algorithm computes probabilities for eg.:
 - $P(I-A \mid I-A, NN)$
 - $P(B-A \mid I-A, NN)$ and even 'emission' probabilities from the $P(NN \mid O)$ (from unobserved to observed)
- Based on the corpus with labelled sequences, the most or least likely ones will then be determined.

B-A: beginning of aspect

I-A: part of aspect

O: anything out of vocabulary



Implicit aspect

- Explicit aspects
 - “*The **picture quality** of the camera is **great**.*”
- Implicit aspects - opinion expressions that do not mention aspect explicitly as a noun
 - “*The camera is **expensive***” (“price”)
 - “*This **beautiful camera**...*” (“appearance”)
 - “*The camera **does not easily fit in a pocket***” (“size”)
- Mapping adjectives or verb phrases to aspects
 - Corpus-based approach: using the co-occurrence of sentiment words and explicit mentions e.g. “*The **size of the phone is small***” vs “*the phone is small*”
 - Dictionary-based approach: using dictionary definition of “***expensive***” – “*marked by high prices*”



Grouping aspects into categories

- People use different words or phrases to describe the same aspect or aspect category.
e.g. “*sound quality*”, “*voice quality*”
- Need to group **aspect expressions** into **aspect categories**, each representing one aspect
- Very challenging as it’s **subjective** task. Different application or different users may require different categories based on application need or granularity of analysis
- Can use WordNet or other thesaurus to find synonyms, but not sufficient
 - Domain dependent synonyms
 - Multiword phrases, not in WordNet or dictionaries
 - Not synonyms – “expensive” and “cheap”



Some methods

- Aspects sharing common words: “*battery life*” and “*battery power*”
- Aspects that are synonyms in dictionaries: “*movie*” and “*film*”
- Aspects with short lexical distances measured in WordNet: “*movie*” and “*show*”
- Mapping aspect expressions to aspect nodes in an existing taxonomy, based on similarity
- Using topic modelling (LDA)
- Other semi-supervised methods



RECENT ADVANCES IN NLP (AND SENTIMENT ANALYSIS)



Pre-2012

- Bag of words approach
 - does not consider the order (semantics nor context) of words:
 - Term Document Matrix, tf-idf indexing, etc.
- Lots of work at document classification, like SVC, MaxEnt, NB, etc.
- Some attempts to get to the semantics, like Latent semantic indexing
- Incremental improvement by **feature engineering** (still useful)



2013 to 2016 – Word Embeddings

- Vectorization using word embedding
 - A vector representation of a word (word vectors) from contextual corpus (map a word using a corpus into a n-dimensional vector)
 - Word2Vec (2013 by Google) or GLOVE (2014 by Stanford)
- Task:
 - Word2vec - using the word itself to predict its probable surrounding words.
 - GLOVE - using the contextual *window* words to predict the word itself.
- Resulting in embeddings that somehow capture the semantics of words



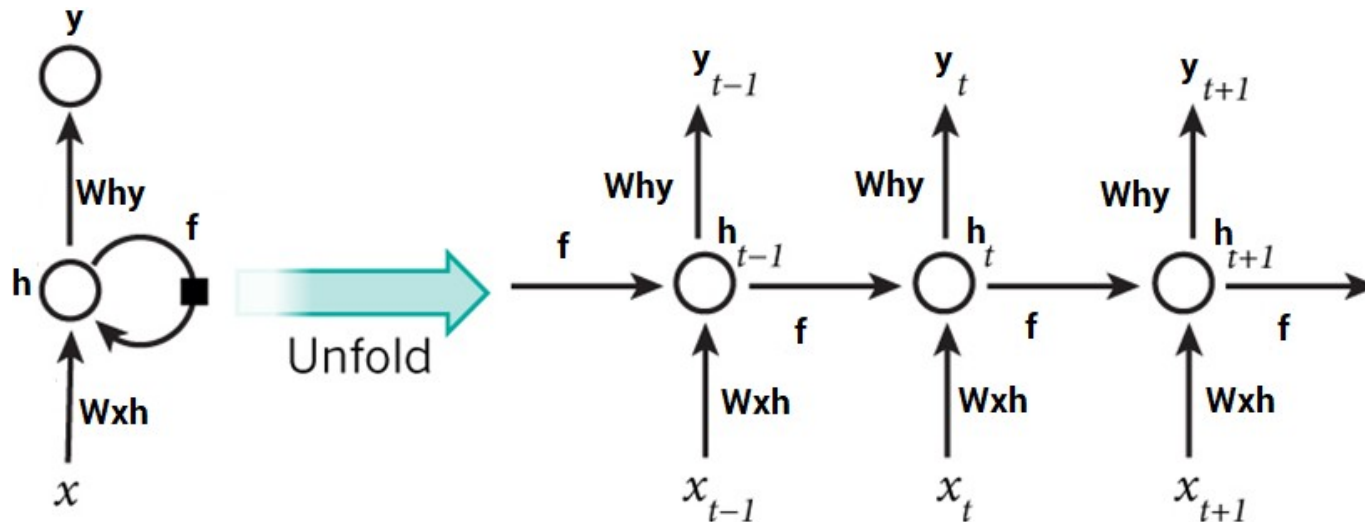
Deep Learning

- The latest techniques in NLP come from Deep Learning.
 - FNN – feedforward neural network
 - CNN – convolutional neural network
 - RNN – recurrent neural network (Sequence modelling!)
- This will be covered in greater detail in the next course - TPML


Recurrent Neural Network

- This is commonly used for NLP.
- Learn through **sequential dependence** of x_t (eg time series data, language sequences)
- Popular improved variant eg. LSTM (long-short term memory)

$$f(x_{1t}, x_{2t}, x_{3t} \dots) \rightarrow y_{1t}, y_{2t}, y_{3t}$$



Long-Short Term Memory



LSTM 'remedies' RNN problems of vanishing gradients (due to the long sequence) by having a sigmoidal 'forget' gate that chooses whether to include the h_{t-n} as either 1 or 0.

- Many NLP tasks can benefit from such sequence modelling. However there are still issues
 - Vanishing gradients
 - Cannot parallelize
- Solution: **Transformer**
 - A novel architecture of ML using attention (focuses on the features that really matter)
 - The ‘transformer’ here is an encoding/ decoding layer with a special attention layer
 - Better handling of long term dependencies

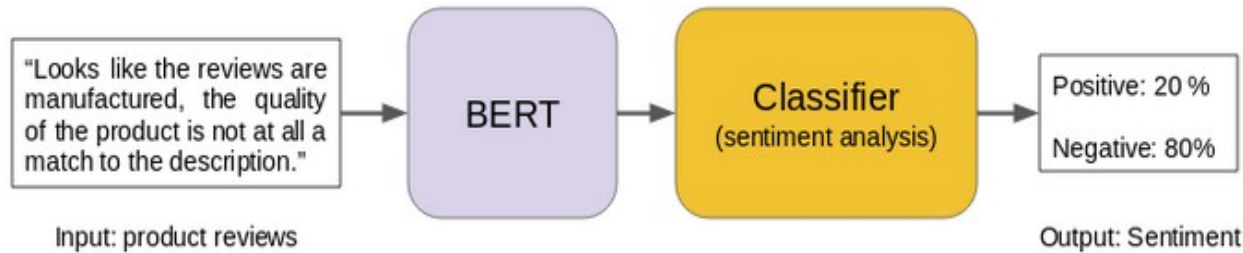


Huge advances from 2018

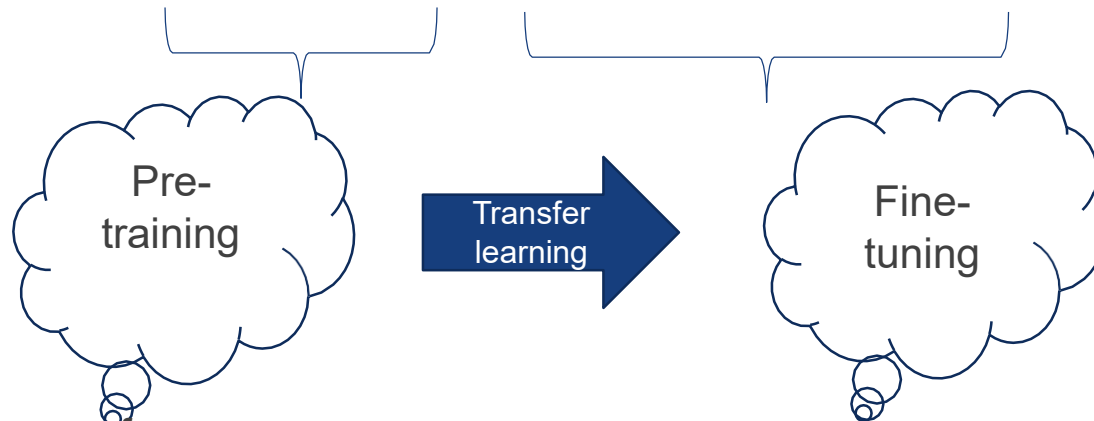
- Elmo (Mar 2018) – contextualized word embeddings
- ***Pre-trained*** models
 - Pre-trained transformers using huge amount of text data
 - The pre-training helps to capture context, syntactic and semantic information in the text.
 - Can be further fine-tuned using small data set to perform a wide variety of NLP tasks, including sentiment mining tasks
 - BERT, GPT/GPT2/GPT3, XLNet, T5, BART, etc.



BERT for sentiment analysis



Fine-Tuning BERT on a Sentiment Analysis Task



By Google –

- BERT base – no of transformer layers = 12; total parameters = 110M
- BERT Large: Number of Transformers layers = 24, Total Parameters = 340M

Sentiment analysis, text classification (eg. spam filters, toxic comments, Q&A)



Leading from the start

Sentence sentiment classification



Rank	Model	Score	CoLA	SST-2	MRPC	STS-B	QQP	MNLI-m	QNLI	RTE
1	BERT: 24-layers, 1024-hidden, 16-heads	80.4	60.5	94.9	85.4/89.3	87.6/86.5	89.3/72.1	86.7	91.1	70.1
2	Singletask Pretrain Transformer	72.8	45.4	91.3	75.7/82.3	82.0/80.0	88.5/70.3	82.1	88.1	56.0
3	BiLSTM+ELMo+Attn	70.5	36.0	90.4	77.9/84.9	75.1/73.3	84.7/64.8	76.4	79.9	56.8

A few years back from
<https://gluebenchmark.com>



A race is going on...

Rank Name		Model	URL	Score	CoLA	SST-2	MRPC	STS-B	QQP	MNLI-m	MNLI-mm	QNLI	RTE	WNLI	
1	AliceMind & DURL	StructBERT + CLEVER		91.0	75.3	97.7	93.9/91.9	93.5/93.1	75.6/90.8	91.7	91.5	97.4	92.5	95.2	
2	ERNIE Team - Baidu	ERNIE		90.9	74.4	97.8	93.9/91.8	93.0/92.6	75.2/90.9	91.9	91.4	97.3	92.0	95.9	
3	DeBERTa Team - Microsoft	DeBERTa / TuringNLRv4		90.8	71.5	97.5	94.0/92.0	92.9/92.6	76.2/90.8	91.9	91.6	99.2	93.2	94.5	
4	HFL IFLYTEK	MacALBERT + DKM		90.7	74.8	97.0	94.5/92.6	92.8/92.6	74.7/90.6	91.3	91.1	97.8	92.0	94.5	
+	5	PING-AN Omni-Sinitic	ALBERT + DAAF + NAS	90.6	73.5	97.2	94.0/92.0	93.0/92.4	76.1/91.0	91.6	91.3	97.5	91.7	94.5	
6	liangzhu ge	Deberta + adv (ensemble)		90.4	72.7	97.3	92.7/90.3	93.2/92.9	75.6/90.8	91.7	91.5	96.4	92.5	95.2	
7	T5 Team - Google	T5		90.3	71.6	97.5	92.8/90.4	93.1/92.8	75.1/90.6	92.2	91.9	96.9	92.8	94.5	
8	Microsoft D365 AI & MSR AI & GATECH	MT-DNN-SMART		89.9	69.5	97.5	93.7/91.6	92.9/92.5	73.9/90.2	91.0	90.8	99.2	89.7	94.5	
+	9	Huawei Noah's Ark Lab	NEZHA-Large	89.8	71.7	97.3	93.3/91.0	92.4/91.9	75.2/90.7	91.5	91.3	96.2	90.3	94.5	
+	10	Zihang Dai	Funnel-Transformer (Ensemble B10-10-10H1024)		89.7	70.5	97.5	93.4/91.2	92.6/92.3	75.4/90.7	91.4	91.1	95.8	90.0	94.5

<https://gluebenchmark.com/leaderboard>




ABSA with DL

- Most treat extraction as a sequence tagging task, using LSTM-based models, or pre-trained transformers.
- An example solution using BART
 - a strong sequence-to-sequence pre-trained model for natural language generation.
 - Encoder and decoder layers
 - Pre-training task: take a masked or permuted sentence as input, return the restored sentence.
- Given input text, generate a sequence of token indexes and class indexes for various ABSA tasks

Yan, Hang, et al. "A Unified Generative Framework for Aspect-Based Sentiment Analysis." *arXiv preprint arXiv:2106.04300* (2021).

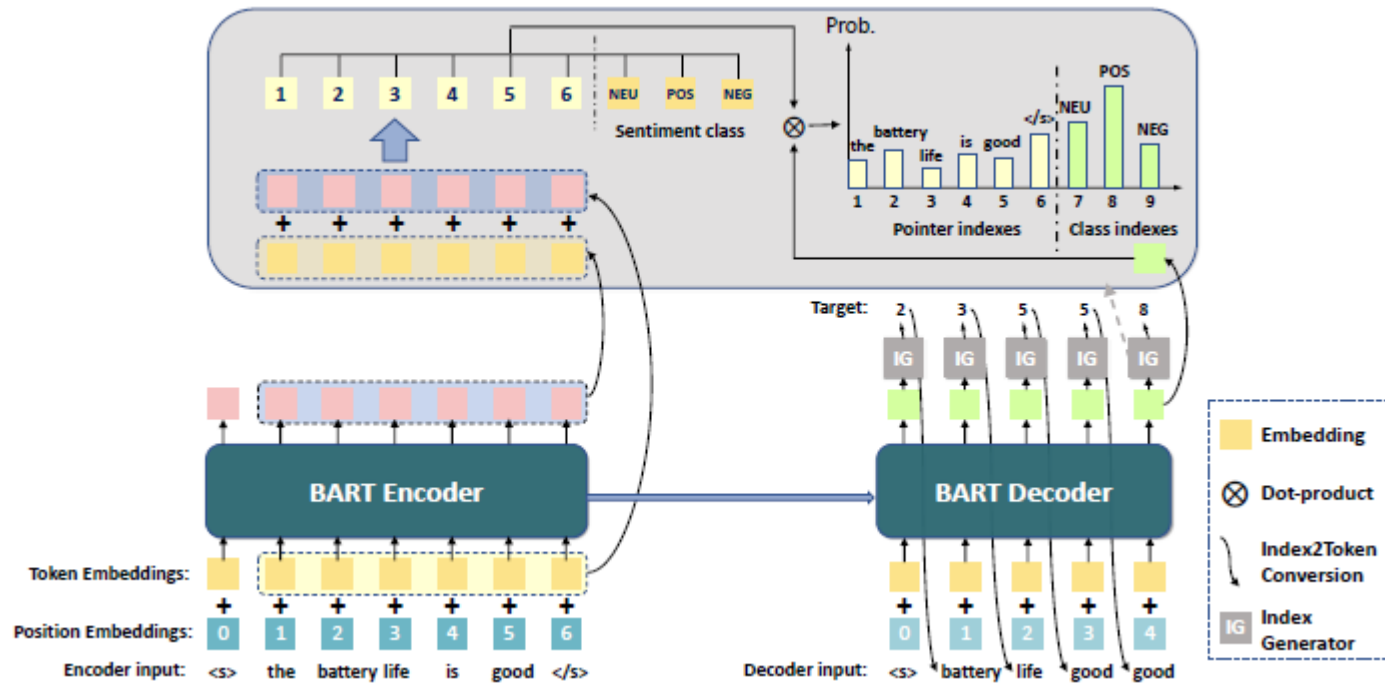
ABSA subtask formulation



Subtask	Target Sequence
<i>AE</i>	1, 2, 12, 12, </s>
<i>OE</i>	4, 4, 7, 8, 14, 14, </s>
<i>ALSC</i>	<u>1</u> , <u>2</u> , POS, </s>
	<u>12</u> , <u>12</u> , POS, </s>
<i>AOE</i>	<u>1</u> , <u>2</u> , 4, 4, 7, 8, </s>
	<u>12</u> , <u>12</u> , 14, 14, </s>
<i>AESC</i>	1, 2, POS, 12, 12, NEG, </s>
<i>Pair</i>	1, 2, 4, 4, 1, 2, 7, 8, 12, 12, 14, 14, </s>
<i>Triplet</i>	1, 2, 4, 4, POS, 1, 2, 7, 8, POS, 12, 12, 14, 14, POS, </s>



Overall architecture



Source sequence: “<s> the battery life is good </s>”

Target sequence: “2 3 5 5 8 6”



References

- Liu, Bing. "**Chapter 6: Aspect and Entity Extraction**". *Sentiment Analysis: Mining Opinions, Sentiments, and Emotions*. Cambridge University Press, 2015.
- Liu, Bing. "Sentiment analysis and opinion mining." *Synthesis Lectures on Human Language Technologies* 5.1 (2012): 1-167.
- Pang, Bo, and Lillian Lee. "Opinion Mining and Sentiment Analysis." *Foundations and Trends in Information Retrieval* 2.1-2 (2008): 1-135.
- Stanford Typed Dependencies Manual
(https://nlp.stanford.edu/software/dependencies_manual.pdf)
- Universal Dependencies
(<https://universaldependencies.org/u/dep/all.html/>)